



РЕШЕНИЯ РЕЧЬ

Speech intelligence that passes the human test.

Синтез и транскрипция речи для естественного, выразительного и тонкого взаимодействия.

Голосовые агенты



Банковский
Агент

Привет, чем я могу
помочь тебе сегодня?

Я хотел бы открыть
бизнес-счет. Что вам
от меня нужно?

Позвольте мне
получить ваш
идентификатор в
файле.

Text-to-speech



Car assistant

Traffic ahead on the A40 near
Hillingdon. I've rerouted you via
the M25. Estimated arrival is
now 9.15am, which still gives

Speech-to-text



Ir

S S

00:01-

The D
billion
sugge

00:08-

We ne
before

S S

00:12-



Обратитесь в отдел продаж ➤



На основе современных открытых голосовых моделей.



Voxtral TTS.

Реалистичная, эмоционально выразительная генерация и клонирование голоса.



Voxtral Mini Transcribe 2.



Voxtral в реальном времени.

Транскрипция прямых трансляций с задержкой менее 200 мс.

Почему речь Мистраля?

Агенты, которые звучат как люди.



Услышь и разберись во всем.



Локализация для разных языков и акцентов.



Работает на основе ваших моделей в вашей инфраструктуре.





Get started ➔

API.

Programmatic access to Mistral's audio models for custom integrations.



Get started ➔

Playground.

Test voice generation, cloning, and transcription in Mistral Studio.



Get started ➔

Enterprise.

Custom deployments, solutions, model training, and dedicated support.



Closing the loop on audio intelligence.

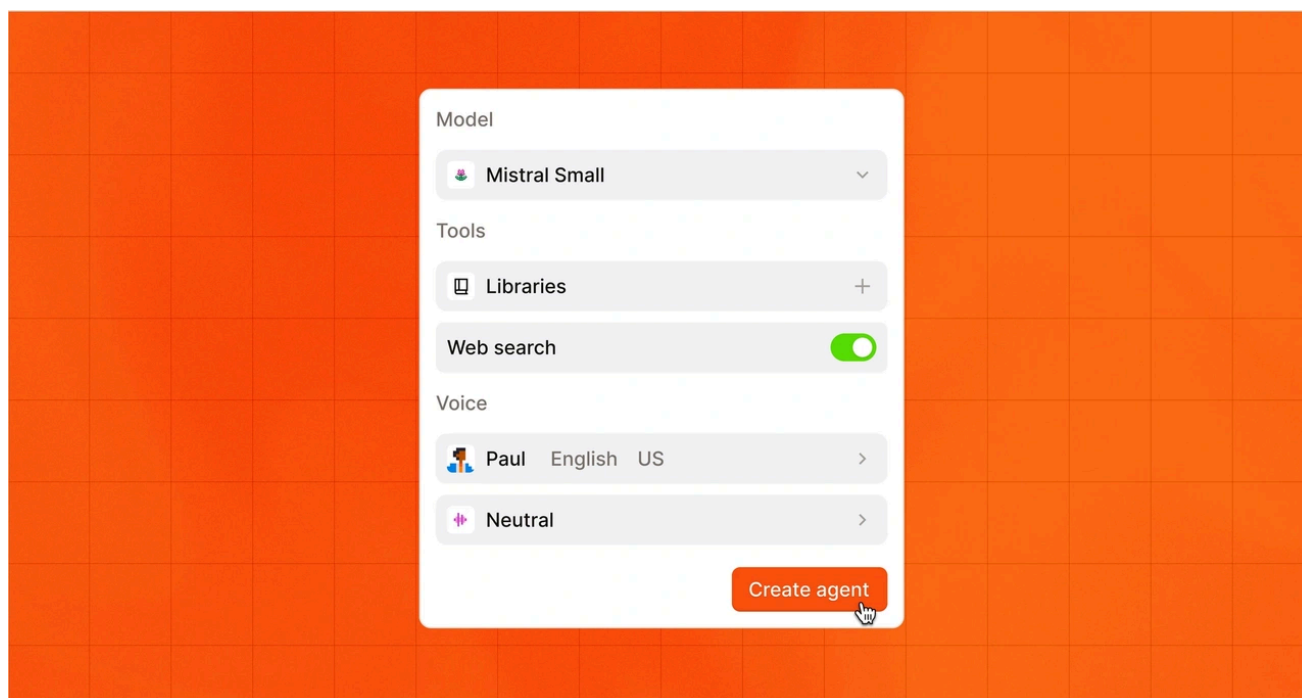
Construct and customize



Construct and customize voices.

Voice agents.

Real-time voice-to-voice conversations that listen, reason, and respond with your brand's voice, tone, and domain knowledge.

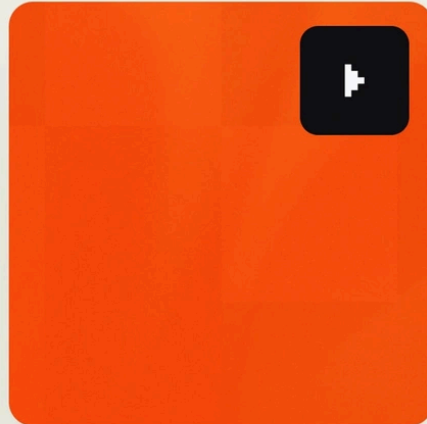


Voice cloning.

Replicate any voice from a sample as short as 3 seconds, capturing tone, rhythm, and personality.



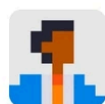
ORIGINAL VOICE



GENERATED VOICE

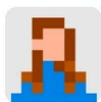
Text-to-speech.

Emotionally expressive speech and voice cloning that captures a speaker's personality. Adapt to any voice from a short sample, or use preset voices.



Paul

Male • English • US



Marie

Female • French



Oliver

Male • English • UK





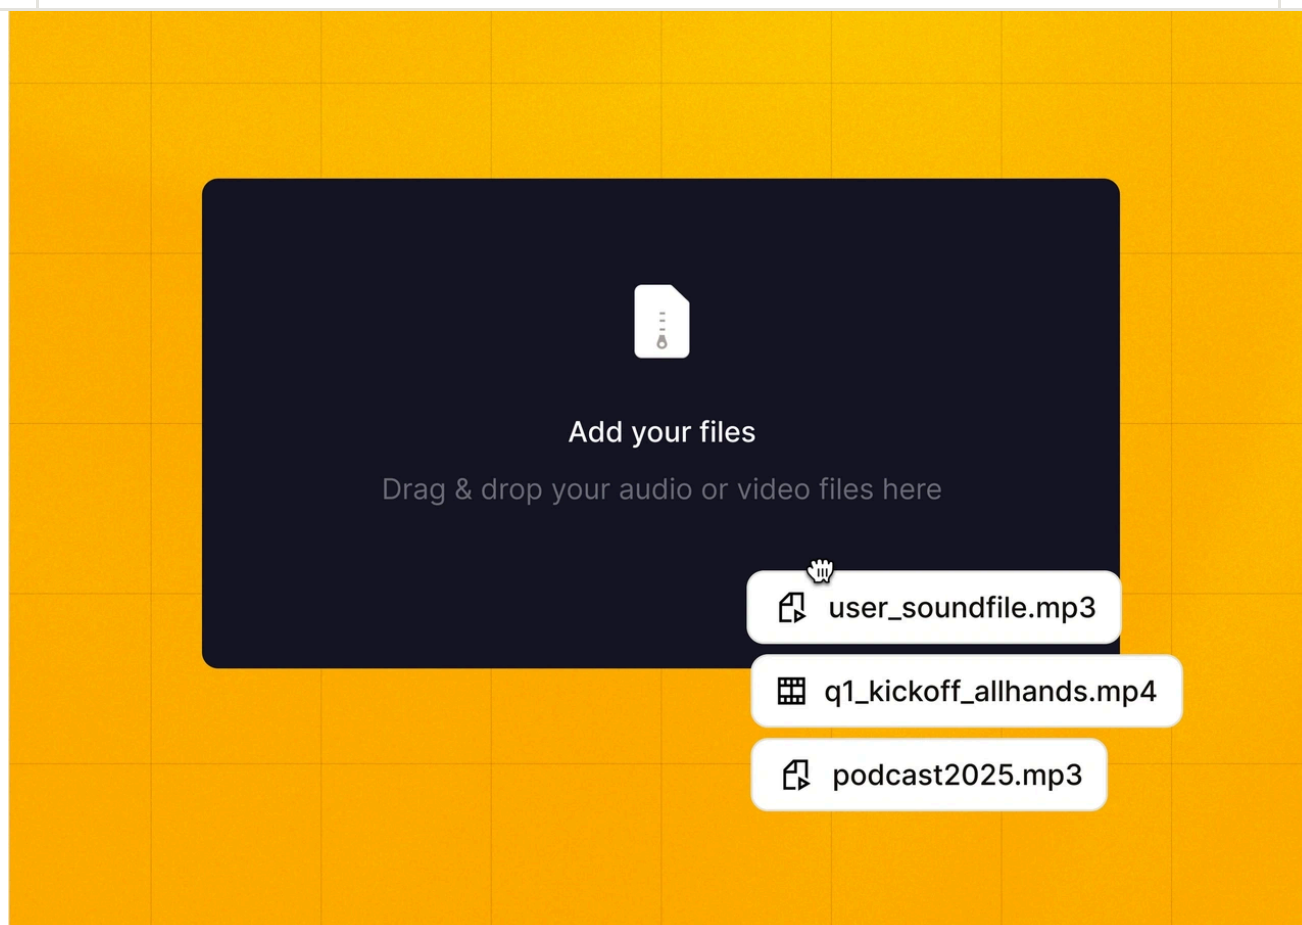
Capture every word.

Real-time transcription.

Streaming architecture that transcribes as audio arrives, not in chunks, with latency configurable down to sub-200ms.



You can see in **real time** what you're dictating




Batch transcription.

Process hours-long meetings, call recordings, and compliance archives, with structured outputs and speaker attribution.

Cross-lingual adaption.

Generate speech in one language using a voice from another, preserving accent and identity.



Prototype, test, tune, adapt.

Audio playground.

Test conversations, voice generation, and transcription in Mistral Studio with actors, voice emulation, diarization, and per-input controls.



Home

API Keys

Build

Playground

Document AI

Audio

Speech to Text

Text To Speech

Voice Agents

Agents

Workflows

Audio Playground

Get started

Generate text

Write a short story

Make an announcement

Begin a conversation

Ask a Technical

Dial a call centre

Start a podcast

Speaker diarization.

Identifies who said what and when, with speaker labels and start/end timestamps for meetings, interviews, and multi-party calls.


0:47 **S** Speaker 1

We need the infrastructure audit done before the migration window opens.

0:58 **S** Speaker 2

The audit is scheduled for next week. Can we align the security review to run in parallel?

Context biasing.



How teams use Mistral Speech today.



Customer support.

Voice agents that route and resolve queries across



Financial services.

Compliant voice AI for wealth management advisory,



Manufacturing industrial operations.

Voice interfaces for inspection, product feedback, and field



Обратитесь в отдел продаж ➤



Resources.



Introduction to audio capabilities ➤

Documentation.



Introducing Voxtral TTS ➤

News.



Follow updates on X ➤

Community.



Frequently asked questions.

How do I get started with Mistral Speech?



Does Mistral have a speech transcription model?



Can I clone Voxtral to my own or my brand's voice?



What languages are supported for transcription and voice generation?



How does the full speech-to-speech pipeline work?



Can I deploy Mistral Speech on-premises or in a VPC?



BUILD YOUR AI FUTURE.

The most expressive, accurate,
and open audio AI for the
enterprise.





Products

Vibe

Vibe Code

Studio

Forge

Compute

Pricing

Solutions

Delivery methodology

Model customization

Coding

Document intelligence

Speech

Mistral for finance

Mistral for public institutions

Mistral for manufacturing

Why Mistral

About us

Careers

Partners

Our customers

Our models

Brand

Legal

Terms of Service

Privacy Policy

Privacy choices

Data processing agreement



Обратитесь в отдел продаж ➔



Get Mistral Vibe

