Привязка к инфраструктуре, которая обходится компаниям, работающим в сфере ИИ, в сотни миллионов долларов

ИИ развивается быстрее, чем аппаратное обеспечение, на котором он работает. Вот почему Nvidia, AMD и гиперскейлеры пересматривают инфраструктуру ИИ, чтобы избежать дорогостоящей привязки.

30 июня 2026 г. в 15:04 Аманда Касвелл

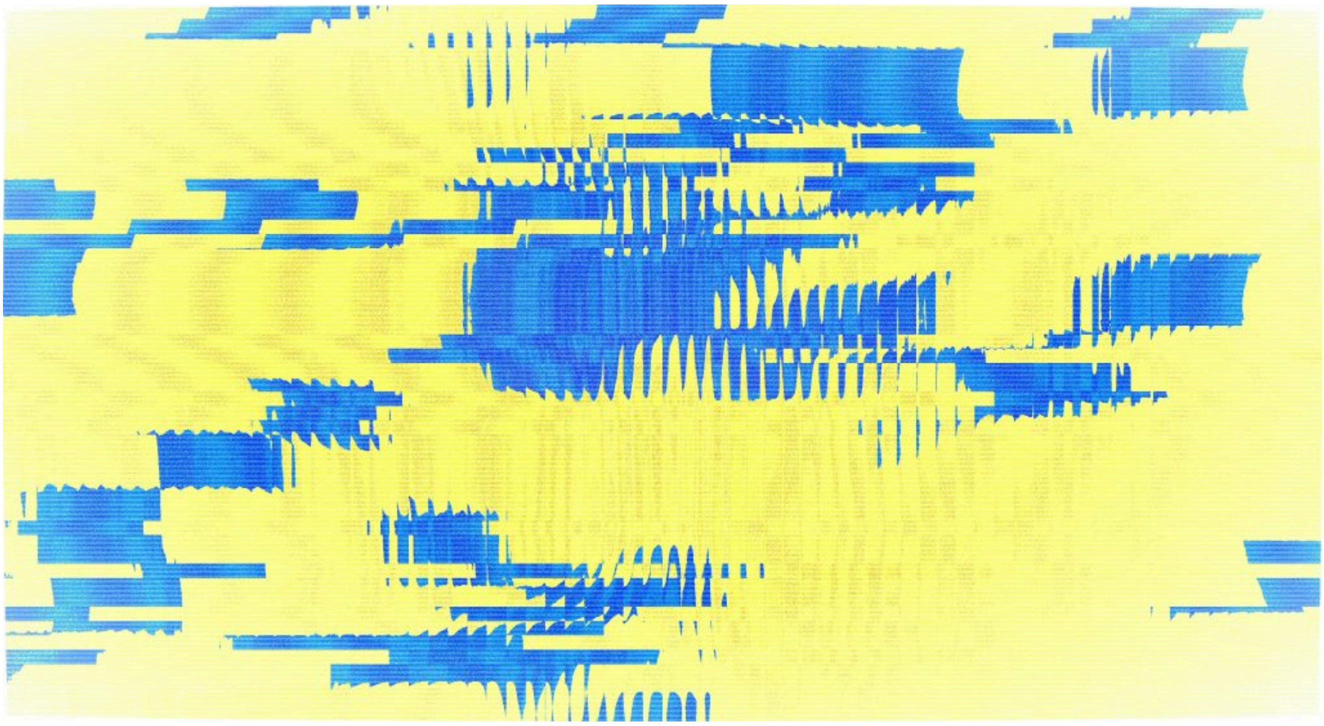

Егор Комаров для Unsplash

В течение двух лет в гонке за инфраструктурой для ИИ доминировал один вопрос: у кого самый быстрый графический процессор? Джим Келлер считает, что это неправильный вопрос.

В недавнем интервью*EE Times* генеральный директор Tenstorrent утверждает, что самый рискованный шаг, который организация может сделать прямо сейчас, — это оптимизация инфраструктуры ИИ для моделей, которые используются сегодня. Дело не в том, что эти модели плохие, а в том, что через 18 месяцев это будут уже не те модели, которые используются сейчас.

пиковой производительности с плавающей запятой.

Не потому, что эти модели плохи, а потому, что через 18 месяцев это будут уже не те модели, которые он использует.

Это похоже на рекламу продукта от Келлера, но за этими словами стоит реальная проблема, потому что искусственный интеллект развивается быстрее, чем инфраструктура, на которой он основан. И компании, потратившие сотни миллионов на разработку одного поколения моделей, теперь сталкиваются с необходимостью делать все заново.

Этот страх называется «привязкой», и он меняет отношение крупнейших игроков в сфере ИИ к оборудованию.

Нагрузка на графические процессоры превысила возможности

В 2023 и 2024 годах инфраструктура искусственного интеллекта представляла собой относительно простую задачу: обучить большие языковые модели, предоставить их пользователям и закупить столько графических процессоров, сколько сможет поставить Nvidia. Нагрузка была предсказуемой, и графические процессоры хорошо с ней справлялись.

Затем искусственный интеллект перерос инфраструктуру, для которой он был создан.

Модели логического вывода тратят больше времени на решение проблем, а не на поиск готового ответа. Агенты переключаются между API, базами данных и кодом, прежде чем выполнить задачу. Мультимодальные модели сочетают текст с изображениями, аудио и видео. Ни одна из этих задач не нагружает оборудование так же сильно, как другие, и это заставляет команды, отвечающие за инфраструктуру, пересмотреть предположения, которые казались вполне логичными всего два года назад.

TRENDING STORIES

2. Palantir and Nvidia want to change who owns government AI
3. Neoclouds, sovereign AI and Postgres: The new operating model for regulated enterprises
4. GitLab just surveyed 1,500 developers. Here's why it matters for your codebase.
5. OpenAI wants to claim more of the AI stack with Jalapeño, its first custom chip

AI outgrew the infrastructure it had been built for.

No single chip architecture handles all of that equally well. And the organizations building AI infrastructure are starting to realize that the question isn't just *which accelerator is fastest* — it's *how do we build systems that won't need to be torn apart every time AI takes another leap?*

Nvidia is already selling one answer

Look at what Jensen Huang has been talking about, and it's not GPUs anymore.

At GTC 2026, Nvidia unveiled the Vera Rubin platform — seven chips designed to operate as a single system: the Rubin GPU, Vera CPU, NVLink 6 switch, ConnectX-9 networking, BlueField-4 DPU, and more. The Vera CPU exists

for the CPU-intensive work of agentic AI — tool calls, code execution, orchestration. Nvidia calls these deployments "AI factories," and the language is deliberate. They're selling complete infrastructure, not individual accelerators.

When the company with 70% market share stops leading with GPU benchmarks and starts talking about system-level co-design, it tells you where the center of gravity in this market is moving.

That reframing matters. When the company with 70% market share stops leading ...h GPU benchmarks and starts talking about system-level co-design, it tells you ...ere the center of gravity in this market is moving. Compute still matters. But

Hyperscalers design their own silicon

AMD sees the same problem, even if it's taking a different route. Helios brings together CPUs, GPUs and networking into one rack-scale platform, reflecting a broader shift away from treating the GPU as the center of the universe. The pitch isn't "our accelerator is faster." It's that the infrastructure surrounding the chip increasingly matters just as much as the chip itself.

The hyperscalers have been making this argument with their wallets for even longer. Google has spent a decade co-designing its TPU silicon, interconnects and software framework — its seventh-generation Ironwood chip is now generally available — giving it unusual control over the full stack. Amazon went the opposite direction, building separate chips for separate jobs: Trainium for training, Inferentia for inference, with Trainium3 now in production and serving customers like Anthropic. Microsoft's Maia 200 targets inference costs, while the company simultaneously deploys Nvidia's Vera Rubin NVL72 for training and experimentation — arguably the most pragmatic dual-track strategy in the market.

Behind all of them sits Broadcom, whose AI semiconductor revenue recently crossed $10 billion in a single quarter, driven by staggering demand for custom accelerators and data center switches. Broadcom designs custom accelerators for Google, Meta and others while supplying the Tomahawk and Jericho switch silicon that connects those accelerators at data center scale. Custom ASIC shipments are projected to grow roughly 45% year over year in 2026 — triple the growth rate of merchant GPUs.

Then there are the companies that decided building a better GPU wasn't the answer.

Cerebras questioned the need for thousands of interconnected chips, opting instead for a wafer-scale processor that keeps far more of the workload on a single piece of silicon. Groq took the opposite approach, optimizing almost entirely for inference. SambaNova focused on enterprise AI, building systems where efficiently serving multiple models matters more than posting the fastest benchmark.

Adaptability beats raw speed

first wave of generative AI rewarded whoever could buy the most compute. ι nat made sense when most organizations were solving the same problem. Today, Keller's example illustrates the point. Tenstorrent's BlackHole architecture uses standard Ethernet instead of proprietary interconnects, allowing its hardware to slot alongside existing GPU deployments rather than replacing them. Keller told EE Times that one customer used Tenstorrent's Galaxy servers to increase token throughput on GPUs they already owned instead of rebuilding their infrastructure from scratch.

Whether Tenstorrent's approach becomes the industry standard is almost beside the point.

The bigger idea is already spreading. Across the industry, companies are spending less time asking how to build the fastest AI hardware and more time asking how to build hardware that won't have to be replaced every time AI takes another leap forward.

The question that matters now

No one knows what AI workloads will look like three or five years from now. That's the problem.

Infrastructure refresh cycles are measured in years. AI models seem to reinvent themselves every few months. Building around today's workloads is starting to look like a risky bet when tomorrow's could demand something different.

> Infrastructure refresh cycles are measured in years. AI models seem to reinvent themselves every few months.

Every company is responding in its own way. They have different strategies but are still asking the same question: **How do you build infrastructure that outlasts the AI running on it?**

That may prove to be a more important engineering challenge than building the next record-breaking accelerator. And it's the challenge driving companies from Nvidia and AMD to Google, Amazon, Broadcom and Tenstorrent.

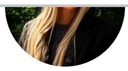

artificial intelligence, developer tools, foundation models, and emerging technologies, with a particular focus...

Read more from Amanda Caswell →

TNS owner Insight Partners is an investor in: Anthropic.

SHARE THIS STORY

TRENDING STORIES

1. AWS, Microsoft, and Google agree the session is the new unit of compute. They disagree on how to isolate it.
2. Palantir and Nvidia want to change who owns government AI
3. Neoclouds, sovereign AI and Postgres: The new operating model for regulated enterprises
4. GitLab just surveyed 1,500 developers. Here's why it matters for your codebase.
5. OpenAI wants to claim more of the AI stack with Jalapeño, its first custom chip

Kong

KAi v2 is Live: The Agent Building Infrastructure for Agents
30 June 2026

Moving from Probabilistic Reasoning to Deterministic Execution
24 June 2026

Kong Mesh 2.14: Finer Zone Proxy Control and Tighter Security
23 June 2026

SerpApi

MCP Apps with FastMCP: Turning Tool Output Into Interactive UI
30 June 2026

How to Scrape Naver AI Briefings with SerpApi
30 June 2026

Connect Sakana AI (Fugu) with Web Search API
30 June 2026

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

FOLLOW TNS

Disclosures
Terms of Use
Advertising Terms & Conditions
Privacy Policy
Cookie Policy