


МОДЕЛИ ИИ / ИСПОЛЬЗОВАНИЕ ИИ / БЕЗОПАСНОСТЬ

Станет ли это Mythos? Мнение одного программиста о сочетании отладки и оптимизации в Anthropic

Сообщество независимых разработчиков тестирует модель искусственного интеллекта Mythos от Anthropic на предмет обнаружения уязвимостей в системе безопасности, проверяя заявления о ее способности выявлять уязвимости.

24 июня 2026 года, 16:51 [Адриан Бриджуотер](#)



Макуде Мариана Куэста

Американские горки от Anthropic с Claude Fable 5 и **Mythos 5** в этом году оставили после себя немало проблем, но **суматоха** не помешала независимым разработчикам тестировать заявленные возможности моделей с помощью собственных методов и каналов.

может находить действительно сложные ошибки в системе безопасности и уязвимости в киберсистемах.

Достаточно богатая смесь для отладки – будет ли это Мифом?

“Идея заключалась в том, чтобы собрать ошибки, которые были специально обнаружены Mythos, как описано в собственной документации **Anthropic** [и начать создавать сервис сравнительного анализа]”, - написал Купер в своем блоге.

«[Я хотел] найти коммит, сделанный до исправления ошибки, убедиться, что модель высшего уровня (**Opus**, в данном случае) может распознать и понять ошибку, если ей прямо указать на нее, и добавить этот коммит в наш корпус для проверки того, могут ли модели, работающие вслепую, точно обнаружить и описать ошибку», — объяснил Купер.

В честь серии рекламных видеороликов **Will It Blend?** на YouTube, которые служили маркетинговым инструментом для линейки блендеров Blendtec, где основатель компании Том Диксон пытался измельчить все подряд — от целой курицы до iPhone 4, — аналитическая статья Купера, опубликованная в мае этого года, получила забавное название Will It Mythos?

В описании его процесса отмечается, что ранее он создал инструмент под названием **Nelson** для автоматизации поиска ошибок в своих проектах. Он отметил, что «уже заметил удивительные различия» между разными моделями (Nelson работает с различными моделями через **Claude Code**, Gemini CLI и API, совместимые с OpenAI) и в том, насколько эффективно они выявляют ошибки.

But he wanted hard numbers, and so he mostly used Claude to cook up a benchmark suite that borrows some code from Nelson and takes Mythos to task.

TRENDING STORIES

1. **Cursor, GitLab and Zed agree GitHub is breaking. They disagree on how to rebuild it.**

-
3. Fable 5 vs Opus 4.8: The real stakes, not the spec sheet
 4. Will it Mythos? One coder's verdict on Anthropic's blend of debugging
 5. Anthropic pauses Claude Agent SDK subscription change on day it was due to take effect

Zero-day vulnerabilities in every major operating system

With his ambitions to test bug gathering according to “Anthropic’s own documentation”, let’s remember here that **Anthropic famously said** that during its testing, it found Mythos Preview was capable of identifying and then exploiting zero-day vulnerabilities in “every major operating system and every major web browser” when directed by a user to do so.

“The toughest bugs are multi-file bugs. The models were free to look at all files, but one often needs to know the context to know that a given usage is a problem. This is a hard problem for any security reviewer, human or AI.”

Ultimately, he wanted to find out whether models going in blind can accurately detect and describe a bug. Cooper’s system was built so that the models used can look at the whole code repository and follow logic across file boundaries, but they’re not told what to look for.

“The toughest bugs are multi-file bugs. The models were free to look at all files, but one often needs to know the context to know that a given usage is a problem. This is a hard problem for any security reviewer, human or AI,” Cooper says.

He said he assumed Mythos has more advanced tooling and perhaps runs the software in a debugger or performs **fuzz testing**.

Some credence that Mythos is particularly good at this

Assessing at everything Mythos might do is beyond the goals of this project for now. But there are bugs in this corpus that are extremely hard to find, giving some

“It almost certainly leads in raw capability, so in that sense, Mythos is engineered to blend an expansive range of bugs. But it’s important to recognize the difference between simply leading a benchmark, versus gating all security progress on one model.” – Conor Sherman, Sysdig.

Raw model capability is not the same as gate-all security

As opinions on Anthropic’s new model abilities start to crystallize, we will gradually gain more real-world insight into their working mechanics.

Conor Sherman, global CISO at runtime security specialist Sysdig is a glass-half-full kind of guy, he tells *The New Stack* that Mythos stands apart from other AI models.

“It almost certainly leads in raw capability,” Sherman says. “So in that sense, Mythos is engineered to ‘blend’ an expansive range of bugs (but perhaps not a whole chicken or iPhone), and it has an inherent ability to perform well when hunting out the multi-file bugs that stump everything else seems to be where it’s most notably ahead.”

But there are caveats here: Sherman advises that it’s important to recognize the difference between simply leading a benchmark and gating all security progress on one model, like some kind of cyber wonder-panacea.

“A less capable but less expensive model wired into the right scaffolding can close most of that gap,” explains Sherman. “The edge that actually matters for defenders isn’t model-versus-model on a corpus of known bugs. Security in an age of AI-powered threats requires runtime context and real-time signals that equip teams to act at the speed at which today’s attackers move.”

No tool should be its own judge

Nhien Renaudineau, co-founder and co-CEO at AI-native synthetic testing company Zark, tells *The New Stack* that in terms of pure capability, agentic testing and

“Systems in this class can already help engineers reason through larger code contexts, identify likely failure paths, and accelerate remediation,” Renaudineau says. “But whatever the benchmark, blend or breed... production reliability is not determined only by what an agent can do.”

Renaudineau draws a line here and says that debugging reliability can only ever be determined by how a tool’s output is verified, and by what we choose to base that verification on – and the first principle here is that an “agent itself cannot be the credible judge” of its own reliability.

“Verification has to be independent of the agent, reproducible and tested against real-world conditions of use – not only against the controlled environment in which the agent was developed or benchmarked,” Renaudineau adds.

What DevSecOps developers should think next

As for the original Will it Mythos man himself, Joe Cooper, noted that on the core question of whether Anthropic’s model can identify really challenging bugs, his benchmark answers “with a resounding, maybe” in this test.

“Mythos, maybe, really is better than the other current models at finding security bugs, as it found four bugs that no model in this experiment found. But I’ll keep testing. It’s possible prompt or tooling or harness changes can enable better results from the current crop of publicly available models,” concluded Cooper. “Over time, I’ll evolve the corpus. It may become a more generic CVE-based benchmark if Anthropic stops bragging about specific bugs.”

With interest and anticipation in Anthropic’s latest batch so fervent and with the work carried out in areas such as **Project Glasswing** being so cloak-and-dagger, the fact that the software engineering community is working to provide so much (hopefully always) impartial analysis must be a good thing.

TNS



Adrian Bridgwater is a technology journalist with three decades of press experience. He has an extensive background in communications, starting in print media, newspapers and also television. Primarily working as an analysis writer dedicated to a software application development ‘beat’,...

TNS owner Insight Partners is an investor in: OpenAI, Anthropic, Sysdig.

SHARE THIS STORY



TRENDING STORIES

1. Cursor, GitLab and Zed agree GitHub is breaking. They disagree on how to rebuild it.
2. Losing Fable made the best case yet for AI models you can run yourself
3. Fable 5 vs Opus 4.8: The real stakes, not the spec sheet
4. Will it Mythos? One coder's verdict on Anthropic's blend of debugging
5. Anthropic pauses Claude Agent SDK subscription change on day it was due to take effect

INSIGHTS FROM OUR SPONSORS



[Deploying Applications to Kubernetes from Source Code: Five Practical Approaches Compared](#)

22 June 2026

[Why Crossplane Alone Is Not Database Management, and What Platform Teams Need Instead](#)

19 June 2026

[Multi-Account AWS Is a Control-Plane Problem — And Most Teams Learn That the Hard Way](#)

13 May 2026

Verification Is Not Testing. Understanding the Difference Could Save Your Codebase

26 June 2026

How long before we stop reading the code?

23 June 2026

The AI Code Verification Bottleneck: Why Faster Code Generation Means Slower Reviews

3 June 2026



Building a Stateful IT Service Desk Agent with LangGraph on Amazon EKS

26 June 2026

Open Governance for MySQL: A Step Forward for the Community

25 June 2026

Governing AI Assets at Scale with MCP Gateway and Registry

17 June 2026

TNS DAILY NEWSLETTER

**Receive a free roundup of the
most recent TNS articles in
your inbox each day.**

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds

THE NEW STACK

About / Contact
Sponsors

 roadmap.sh

Community created roadmaps, articles,
resources and journeys for developers

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

FOLLOW TNS



© The New Stack 2026

[Disclosures](#)

[Terms of Use](#)

[Advertising Terms & Conditions](#)

[Privacy Policy](#)

[Cookie Policy](#)