



Вселенная

Машинное обучение для бизнес-аналитики. 2-е издание

Авторы: Галит Шмуэли, Питер К. Брюс, Питер Гедек, Нитин Р. Патель

Тип файла: pdf

Размер: 43 МБ

Язык: Английский

Страниц: 1188

Машинное обучение для бизнес-аналитики. 2-е издание: концепции, методы и приложения на Python



Введение

Современные компании генерируют огромные объемы данных с веб-сайтов, из мобильных приложений, транзакций клиентов, датчиков, финансовых систем, маркетинговых кампаний и корпоративного программного обеспечения.

Превращение этих данных в полезную информацию для принятия решений — одна из главных целей **бизнес-аналитики**.


Машинное обучение (Machine Learning, ML) стало важной частью этого процесса, поскольку оно позволяет организациям выявлять закономерности, прогнозировать будущие результаты, сегментировать клиентов, обнаруживать необычное поведение и автоматизировать аналитические задачи. Python предоставляет доступ к этим возможностям через такие библиотеки, как Pandas, NumPy, Matplotlib, Seaborn, Scikit-learn и другие специализированные инструменты. 📊


В отличие от традиционной аналитики, которая часто фокусируется на описании произошедшего, машинное обучение может помочь ответить на такие вопросы, как:

- Какие клиенты могут уйти?
- На какие товары может вырасти спрос в следующем месяце?
- Какие транзакции выглядят подозрительно?
- Какие маркетинговые лиды наиболее ценны?
- Как компания может прогнозировать продажи?

- Какие факторы влияют на удовлетворенность клиентов?









Сочетание **бизнес-знаний + данных + машинного обучения + Python** создает мощный аналитический рабочий процесс для студентов, аналитиков, инженеров, менеджеров и специалистов по работе с данными.

Теоретическая база

Традиционная бизнес-аналитика обычно опирается на исторические данные и заранее заданные правила. Система отчетности может сообщить менеджеру, что в прошлом квартале продажи выросли, а информационная панель — показать, какой регион принес наибольший доход.

Машинное обучение привносит еще один аспект: **выявление закономерностей на основе существующих данных и использование этих закономерностей для прогнозирования или принятия решений.**

Общий процесс можно представить следующим образом:

Бизнес-задача → Данные → Подготовка → Модель → Оценка → Бизнес-решение

Машинное обучение не заменяет бизнес-аналитику. Напротив, оно расширяет возможности аналитики, позволяя организациям работать со сложными взаимосвязями, которые может быть трудно выявить с помощью правил, написанных вручную.

Описательная, диагностическая, предиктивная и предписывающая аналитика

Бизнес-аналитику часто делят на четыре основные категории:

Описательная аналитика объясняет, что произошло.

Диагностическая аналитика исследует причины произошедшего.

Предиктивная аналитика прогнозирует, что может произойти дальше.

Предписывающая аналитика помогает определить, какие действия могут привести к желаемому результату.

Машинное обучение особенно ценно для предиктивной и предписывающей аналитики, хотя оно может поддерживать и описательные, и диагностические процессы.

Обучение с учителем

При обучении с учителем модель обучается на примерах, где желаемый результат уже известен.

Например, у компании могут быть исторические данные о клиентах, помеченные следующим образом:

- Остался
- Отменен

Модель классификации может учиться на этих примерах и оценивать вероятность того, что новый клиент может отменить подписку.

Обучение без учителя

Обучение без учителя работает с данными, в которых нет заранее определенной цели.



Классический пример — сегментация клиентов. Компания может предоставить алгоритму кластеризации информацию о покупательском поведении, частоте покупок, структуре расходов и вовлеченности.

Алгоритм может выделять группы со схожими характеристиками.

Разработка признаков

Признаки — это переменные, используемые в модели машинного обучения.

Примеры:

- Customer age range
- Number of purchases
- Average order value
- Website visits
- Days since last purchase
- Marketing interaction frequency
- Product category
- Geographic region



Feature engineering transforms raw business data into meaningful variables that algorithms can use effectively. 🧠

Definition

Machine Learning for Business Analytics is the application of machine learning methods to business data to discover patterns, generate predictions, classify observations, identify groups, detect anomalies, and support better organizational decisions.


Python is particularly suitable because it provides a large ecosystem for:


- Data cleaning
- Statistical analysis
- Visualization
- Machine learning
- Automation
- Model evaluation
- Reporting
- Deployment

The popular **Scikit-learn** library provides many algorithms for classification, regression, clustering, preprocessing, and model evaluation.

Step-by-Step Machine Learning Workflow in Python

A successful ML project begins with the business question rather than the algorithm.







Step 1: Define the Business Problem

Start with a measurable objective.

For example:

Predict which customers are most likely to stop using a subscription service.

The objective should connect directly to a business decision.

Step 2: Collect the Data

Potential sources include:

- CRM systems
- ERP platforms
- Transaction databases
- Website analytics
- Customer surveys
- Marketing platforms
- IoT devices
- Financial systems

Python can connect to many data sources and load datasets into Pandas DataFrames.

Step 3: Explore the Dataset

Data exploration helps identify:

- Missing values
- Duplicate records
- Unusual observations
- Incorrect data types
- Imbalanced categories
- Relationships between variables



Visualization is particularly useful at this stage.

Step 4: Prepare the Data

Business datasets rarely arrive ready for machine learning.

Common preparation activities include:

- Handling missing values
- Removing duplicates
- Encoding categorical variables
- Scaling numerical variables when appropriate
- Removing or investigating abnormal records
- Selecting useful features

Step 5: Split the Data

The dataset is generally divided into training and testing portions.

The training data helps the model learn patterns, while the testing data provides an independent evaluation of how well the model performs on unseen observations.

Step 6: Select a Model

The appropriate algorithm depends on the problem.

For example:

- Logistic Regression → classification
- Linear Regression → numerical prediction
- Decision Trees → interpretable prediction
- Random Forest → robust classification or regression
- Gradient Boosting → powerful predictive modeling
- K-Means → customer segmentation

Step 7: Train the Model

Python and Scikit-learn can simplify model training.

A typical workflow involves importing the algorithm, preparing the training dataset, creating the model, and fitting it to the available data.



Step 8: Evaluate Performance

The evaluation metric should match the business objective.

Classification may use:

- Accuracy
- Precision
- Recall
- F1-score
- ROC-AUC

Regression may use:

- Mean Absolute Error
- Mean Squared Error
- R^2

The most important metric is not always the most mathematically popular one. Business consequences should guide metric selection.

Step 9: Interpret the Results

A model may produce excellent predictions but still be unsuitable if decision-makers cannot understand how it should be used.

Explainability becomes particularly important in areas such as finance, insurance, recruitment, and customer management.

Step 10: Deploy and Monitor

After deployment, the model should be monitored.

Customer behavior, market conditions, product offerings, and economic conditions can change. A model that worked well six months ago may gradually lose accuracy.


Comparison of Machine Learning Techniques


Technique	Typical Business Use	Main Advantage	Potential Limitation
Linear Regression	Sales or demand estimation	Simple and interpretable	May struggle with complex patterns
Logistic Regression	Customer churn classification	Easy to interpret	Limited nonlinear behavior
Decision Tree	Risk or customer classification	Highly interpretable	Can overfit
Random Forest	Classification and prediction	Strong general-purpose performance	Less transparent than a single tree
Gradient Boosting	High-performance prediction	Handles complex relationships	Requires careful tuning
K-Means	Customer segmentation	Simple clustering approach	Requires suitable cluster assumptions
Isolation Forest	Anomaly detection	Useful for unusual observations	Results require business interpretation




Diagrams and Tables


A practical business analytics architecture can be represented as:














Business Question	ML Approach	Example Output
Will a customer leave?	Classification	Churn probability
What might sales look like?	Regression/time-series methods	Sales forecast
Which customers behave similarly?	Clustering	Customer groups
Is this transaction unusual?	Anomaly detection	Risk alert
Which leads deserve attention?	Classification/ranking	Lead priority

Examples

Customer Churn Prediction

A telecommunications company can analyze customer tenure, service usage, support interactions, subscription type, and payment behavior.

Python can prepare the historical data and train a classification model.

The resulting system could identify customers who have characteristics associated with previous cancellations.

The marketing team could then prioritize those customers for retention campaigns.

Sales Prediction

A retailer can combine historical sales, product categories, seasonal information, promotions, and regional data.



A predictive model can learn relationships between these variables and historical demand.

Managers can use the resulting forecasts to improve inventory planning.

Customer Segmentation

An online store may have thousands of customers with different purchasing behaviors.

Clustering can help identify groups such as:

- Frequent buyers
- High-value customers
- Occasional shoppers
- Discount-oriented customers
- Inactive customers

Each group can receive a different marketing strategy.

Fraud Detection

Financial organizations can analyze transaction characteristics to identify unusual activity.

Anomaly detection can flag transactions that differ significantly from normal behavioral patterns.

The system does not necessarily declare that every unusual transaction is fraudulent. Instead, it can prioritize transactions for investigation. 🔍

Real-World Applications

Machine learning for business analytics has applications across many industries.

Retail

Retailers use ML for:

- Demand forecasting
- Customer segmentation
- Product recommendations
- Inventory planning
- Promotion analysis
- Fraud detection



Banking and Finance

Financial institutions can use machine learning for:

- Credit risk analysis
- Fraud detection
- Customer segmentation
- Transaction monitoring
- Financial forecasting

Manufacturing

Manufacturers can combine machine learning with operational data for predictive maintenance.

Sensors can provide information about temperature, vibration, operating conditions, and equipment performance.

ML models can help identify patterns associated with potential equipment problems.

Healthcare Business Operations

Healthcare organizations can use analytics to improve:

- Resource allocation
- Appointment planning
- Operational forecasting
- Patient-flow analysis
- Administrative processes

Applications involving sensitive personal data require strong privacy, security, governance, and regulatory controls.

Marketing

Marketing teams can use machine learning to analyze:

- Campaign performance
- Customer engagement
- Lead quality
- Conversion behavior
- Audience segments



This can help businesses allocate marketing resources more efficiently.

Common Mistakes

Starting With the Algorithm

One of the biggest mistakes is selecting a sophisticated model before clearly defining the business problem.

The algorithm should serve the objective—not the other way around.

Ignoring Data Quality

A sophisticated model cannot compensate for unreliable input data.

Incorrect customer records, duplicates, missing information, and inconsistent categories can produce misleading results.

Data Leakage

Data leakage occurs when information unavailable at prediction time accidentally enters the training process.

This can make a model appear exceptionally accurate during testing while performing poorly in real-world use.

Overfitting

An overfitted model learns the training data too closely.

It may perform impressively on historical examples but fail when presented with new business cases.

Using Accuracy Alone

Suppose fraudulent transactions are extremely rare.

A model that predicts “not fraud” for almost every transaction might achieve high accuracy while being practically useless.

The evaluation strategy must reflect business priorities.



Ignoring Model Drift

Business environments change.

Customer preferences, competitors, regulations, economic conditions, and marketing strategies can all affect model performance.

Challenges and Solutions

Challenge	Practical Solution
Missing data	Investigate the reason for missingness and use appropriate preprocessing
Large datasets	Use efficient data processing and sampling strategies
Poor feature quality	Apply feature engineering and domain expertise
Imbalanced classes	Consider suitable metrics and balancing techniques
Overfitting	Use validation, regularization, and appropriate model complexity

Challenge	Practical Solution
Difficult interpretation	Use interpretable models and explainability techniques
Changing behavior	Continuously monitor model performance
Privacy concerns	Apply data governance, access control, and privacy practices

Case Study: Predicting Customer Churn

Consider a fictional subscription-based software company serving professional customers across North America and Europe.

The company notices that customer cancellations have increased.

Management wants an early-warning system.



Business Objective

The objective is to identify customers who may be at increased risk of cancellation so that the customer-success team can intervene.

Data Collection

The analytics team collects:

- Subscription information
- Login activity
- Product usage
- Customer-support interactions
- Account age
- Billing history
- Service plan
- Previous cancellations

Data Preparation

The team removes duplicate records, investigates missing values, converts categorical information into machine-readable features, and creates behavioral indicators.

Model Development

Several classification approaches are evaluated.

The team does not automatically choose the most complicated model. Instead, it considers predictive performance, interpretability, maintenance requirements, and operational usefulness.

Business Integration

Customers identified as higher-risk are displayed in a customer-success dashboard.

Employees can then review the account and determine whether an intervention is appropriate.

Outcome

The most valuable result is not simply a model score.

The real objective is to help the company **identify customers earlier and take useful action**.

This illustrates an important principle:

A successful business ML project connects predictions to decisions.



Essential Tips

Combine Technical and Business Knowledge

Python skills are valuable, but understanding the business process is equally important.

A data scientist who understands how sales teams, engineers, financial analysts, or operations managers actually work can design much more useful models.

Start With Simple Models

A simple baseline provides an important reference point.

If a complicated model offers little practical improvement, the simpler solution may be preferable because it is easier to explain and maintain.

Visualize Before Modeling

Use charts to understand distributions, relationships, outliers, and trends before training a model. 📈

Separate Training and Testing Properly

Always evaluate the model on data that represents observations it did not use for learning.

Focus on Business Impact

Ask:



Will this prediction change a decision?

If the answer is no, the model may not solve a meaningful business problem.

Build Reproducible Python Workflows

Organize notebooks, scripts, datasets, preprocessing steps, and model configurations carefully.

Reproducibility becomes increasingly important when a prototype turns into a production system.

Monitor After Deployment

Track model quality, data quality, prediction distributions, and business outcomes.

A deployed ML model should be treated as a continuously evolving analytical system rather than a one-time project.

Frequently Asked Questions

What is machine learning in [business analytics](#)?

Machine learning in business analytics involves using algorithms to learn patterns from business data and generate predictions, classifications, segments, or other insights that support decision-making.

Why is Python popular for business analytics?

Python provides a large ecosystem for data processing, visualization, statistics, machine learning, automation, and deployment. Libraries such as Pandas and Scikit-learn make many analytical workflows accessible.

Which Python libraries should beginners learn?

A useful starting combination is **Pandas, NumPy, Matplotlib, and Scikit-learn**. As projects become more advanced, additional libraries can be introduced according to the business problem.

Is machine learning better than traditional business analytics?



Not necessarily. Traditional analytics remains extremely useful for reporting, monitoring, and understanding historical performance. Machine learning is an additional tool that becomes valuable when prediction, classification, segmentation, or automated pattern recognition is required.

Which machine learning algorithm is best for business analytics?

There is no universally best algorithm. The appropriate method depends on the problem, data quality, interpretability requirements, prediction objective, computational resources, and business consequences.

Can machine learning predict sales?

Yes. Machine learning can be used to support sales and demand forecasting when suitable historical and contextual data are available. However, forecasts should account for changing market conditions and should be regularly evaluated.

How important is data cleaning?

Extremely important. Poor-quality data can significantly reduce model reliability. Data preparation is often one of the largest components of a practical machine learning project.

Do business analysts need advanced mathematics to use machine learning?

A strong understanding of basic statistics and analytical concepts is helpful. Beginners can build useful models using high-level Python libraries without manually implementing every mathematical operation. However, advanced professionals benefit from deeper knowledge of statistics, optimization, probability, and model theory.

Conclusion

Machine learning has transformed business analytics from a primarily historical discipline into a more predictive and intelligent decision-support capability. 🚀

Using Python, students and professionals can build workflows that transform raw business data into useful predictions, customer segments, anomaly alerts, and operational insights.



The essential process is straightforward:

Define the business problem → collect reliable data → prepare the dataset → engineer useful features → select an appropriate model → evaluate it properly → interpret the results → connect predictions to decisions → monitor performance.

The most successful machine learning projects are not necessarily those using the most complicated algorithms. They are the projects that solve a genuine business problem, use trustworthy data, produce understandable results, and lead to measurable improvements.

For engineers, analysts, developers, managers, and students, learning **machine learning + Python + business analytics** provides a powerful combination for working with modern data-driven organizations.

Ultimately, the goal is not simply to build a machine learning model. The goal is to transform data into **better decisions, smarter processes, and measurable business value.** 📊🤖👤

Related Posts:



Preview PDF Book



[← Previous Book](#)

[Next Book →](#)

EngiVerse

Explore a vast library of free engineering ebooks curated to help you expand your technical knowledge, conduct groundbreaking research, and advance your career. Our mission is to create a hub of knowledge that fosters innovation, empowers aspiring engineers, and bridges the gap between education and industry.

[Home](#) [Books](#) [About Us](#) [Contact](#)
[Privacy Policy](#) [DMCA](#)

copyright 2025 EngiVerse. All Right Reserved

[Privacy and cookie settings](#)

Managed by Google. Complies with IAB TCF. CMP ID: 300

