7 мая 2026 года Продукт Релиз

Развитие голосового интеллекта с помощью новых моделей в API

Новое поколение голосовых моделей, которые могут анализировать, переводить и расшифровывать речь в режиме реального времени.

Прослушать статью 8:31

Мы представляем три аудиомодели в API, которые открывают перед разработчиками новые возможности для создания голосовых приложений. С помощью этих моделей разработчики могут создавать голосовые интерфейсы, которые звучат более естественно,

- **GPT-Realtime-2**, наша первая голосовая модель с логическим мышлением на уровне GPT-5, которая может обрабатывать более сложные запросы и вести диалог в естественном ключе.
- **GPT-Realtime-Translate**, новая модель для перевода в режиме реального времени, которая переводит речь с более чем 70 языков на 13 языков, не отставая от говорящего.
- **GPT-Realtime-Whisper**, новая модель преобразования речи в текст в режиме реального времени, которая расшифровывает речь по мере ее произнесения.

Попробуйте GPT-Realtime-2

Запустите сеанс и общайтесь с GPT-Realtime-2 в естественном режиме.

Начало сеанса

О чем я могу спросить?

Эта демонстрация ограничена по времени. Используя ее, вы соглашаетесь с Условиями OpenAI и подтверждаете, что ознакомились с нашей Политикой конфиденциальности.

самых естественных способов использования программного обеспечения. С его помощью можно попросить о помощи за рулем, изменить маршрут во время прогулки по аэропорту, получить поддержку на предпочитаемом языке или выполнить задачу, не отвлекаясь на набор текста.

Но для создания полезных голосовых продуктов недостаточно просто быстро реагировать на реплики и воспроизводить естественный голос. Голосовой помощник должен понимать, что имеет в виду собеседник, учитывать контекст, восстанавливаться после изменения запроса, использовать инструменты во время разговора и отвечать так, чтобы это соответствовало ситуации.

Вместе с тем модели, которые мы внедряем, превращают аудио в режиме реального времени из простого диалога с голосовым управлением в голосовой интерфейс, который действительно может выполнять работу: слушать, анализировать, переводить, расшифровывать и действовать по ходу разговора.

Голос как интерфейс взаимодействия между людьми и продуктами

As voice becomes a more natural way to use software, we’re seeing developers build around three emerging patterns in voice AI:

- **Voice-to-action,** where people can describe what they need and the system can reason through the request, use tools, and complete the task. For example, Zillow is building an assistant that can listen, reason, and act on tour for Saturday."

- **Systems-to-voice,** where software can turn context into live spoken guidance. For example, a travel app could proactively tell a traveler: "Your inbound flight is delayed, but you can still make your connection. I found the new gate, mapped the fastest route through the terminal, and your bag is still expected to transfer."
- **Voice-to-voice,** where AI can help live conversations continue across languages, tasks, or changing context. For example, Deutsche Telekom is building voice support experiences where customers can speak in the language they're most comfortable using, while the model translates the conversation in real time.

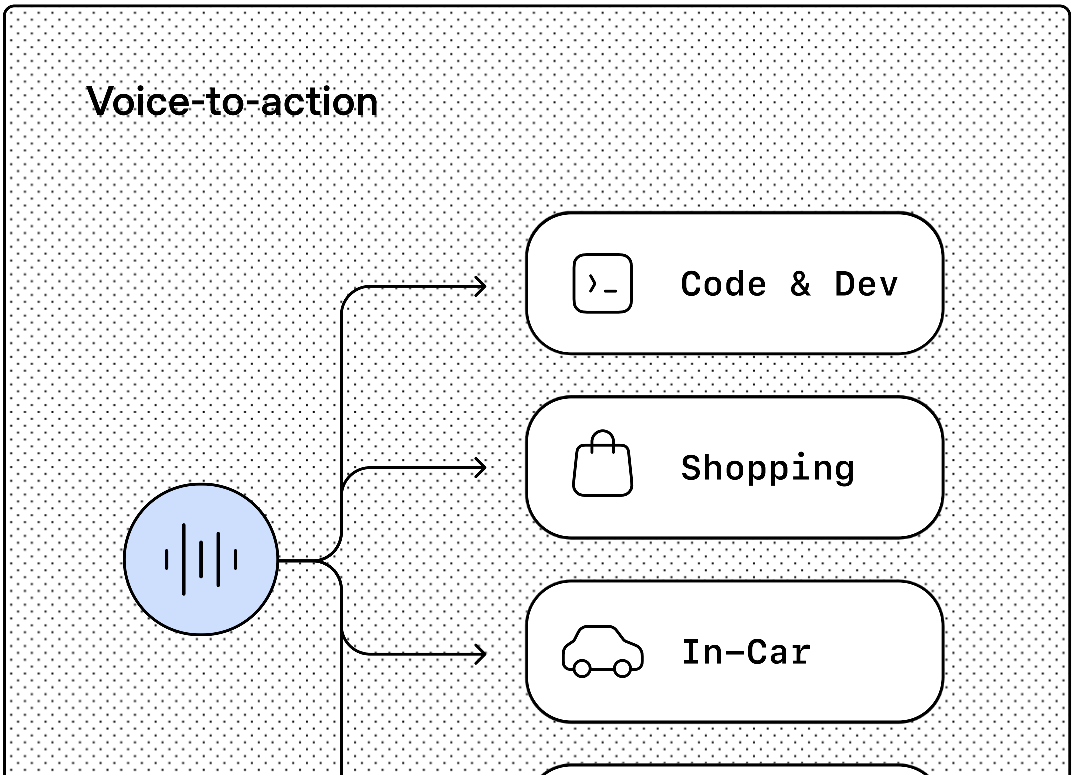

These patterns can also work together. Priceline is working toward a future where travelers can manage entire trips by voice: searching for flights and hotels conversationally, handling changes like adjusting a hotel reservation after a flight delay or getting real-time updates on TSA wait times, and translating conversations once travelers are on the ground.

Realtime voice: helping voice models reason and take action

GPT-Realtime-2 is built for live voice interactions where the model keeps the conversation moving while it reasons through a request, calls tools, handles corrections or interruptions, and responds in a way that fits the moment.

- **Preambles:** Developers can enable short phrases before a main response, like "let me check that" or "one moment while I look into it," so users know the agent is working on the request.
- **Parallel tool calls and tool transparency:** The model can call multiple tools at once and make those actions audible with phrases like "checking your calendar" or "looking that up now," helping agents stay responsive while completing tasks.
- **Stronger recovery behavior:** The model can recover more gracefully by saying things like "I'm having trouble with that right now," instead of failing silently or breaking the conversation.
- **Longer context for agentic workflows:** We're increasing the context window from 32K to 128K to support longer, more coherent sessions and more complex task flows.
- **Stronger domain understanding:** The model better retains specialized terminology, proper nouns, healthcare terms, and other vocabulary that matters in production settings.

calmly while resolving an issue, empathetically when a user is frustrated, or upbeat when confirming a successful action.

- **Adjustable reasoning effort:** Developers can now select from **minimal, low, medium, high, and xhigh** reasoning levels, with **low as the default**, balancing lower latency for straightforward interactions with more deliberate reasoning for complex requests.

The gains show up on audio evals that map closely to production voice agents: GPT-Realtime-2 (high) scores 15.2% higher on Big Bench Audio for audio intelligence than GPT-Realtime-1.5. GPT-Realtime-2 (xhigh) scores 13.8% higher on Audio MultiChallenge for instruction following, improving over GPT-Realtime-1.5 and showing stronger reasoning, context management, and control in live conversations.

Big Bench Audio evaluates challenging reasoning capabilities in language models that support audio input. Audio MultiChallenge evaluates multi-turn conversational intelligence in spoken dialogue systems, including

The magic of GPT-Realtime-2 shows up across a variety of different use cases:

Strategic reasoning | Tone and expressiveness | Spatial reasoning | Alphanumerics | Logic puzzle

User

I'm considering a 900-square-foot indie coffee shop beside a commuter rail station. Foot traffic peaks Tuesday through Thursday from 7 to 10 a.m.; Mondays, Fridays, and afternoons are much softer. The lease is expensive, but I love the idea of cozy seating, slow pour-overs, and local pastries. Give me a strategic pre-mortem: if this fails after a year, what probably happened? Then suggest the smallest version of the business I should test before committing to the full cafe.

GPT-Realtime-2

0:00 1:04

Transcript

GPT-Realtime-1.5

0:00 0:51

Transcript

During early testing, businesses used GPT-Realtime-2 to build voice agents that help customers and employees get things done through natural conversation:

Zillow Glean Genspark Bluejay Intercom Priceline Foundation Health

"What stood out about GPT-Realtime-2 was the intelligence and tool-calling reliability it brings to complex voice interactions. On our hardest adversarial benchmark, this translates to a 26-point lift in call success rate after prompt optimization (95% vs. 69%). GPT-Realtime-2 is also materially more robust on Fair Housing compliance, which is critical for our business. The combination of agentic competence and guardrail strength is what makes it viable for production voice at Zillow."

— Josh Weisberg, SVP and Head of AI at Zillow

Realtime translation: build live multilingual voice experiences

GPT-Realtime-Translate helps developers build live multilingual voice experiences where each person can speak in their preferred language and hear the conversation translated in real time and read the real time transcriptions. It supports more than 70 input languages and 13 output languages, making it useful for customer support, cross-border sales, education, events, media, and creator platforms serving global audiences.

meaning while keeping pace with the speaker, even when people speak naturally, switch context, or use regional pronunciation and domain-specific language. For example, Deutsche Telekom is testing the model for multilingual voice interactions, where lower latency and stronger fluency can make cross-language conversations feel more natural.

In this video, Vimeo shows how GPT-Realtime-Translate can translate a product education video live as it plays, so global customers can hear updates in their preferred language without waiting for a separately produced version.

BolnaAI Vimeo Deutsche Telekom

"Building voice AI for India means handling diverse regional phonetics. In our

Realtime-Translate delivered 12.5% lower Word Error Rates than any other model we tested, along with lower fallback rates, higher task completion, and latency that sustained natural conversation. It sets a new standard for multilingual voice AI."

— Prateek Sachan, Co-founder & CTO at BolnaAI

Realtime transcription: build low-latency transcription experiences

GPT-Realtime-Whisper is a new streaming transcription model built for low-latency speech-to-text. It transcribes audio as people speak, so live products can feel faster, more responsive, and more natural—from captions that appear in the moment, to meeting notes that keep up with the conversation.

The model makes live speech usable inside business workflows as it happens. Teams can power captions for meetings, classrooms, broadcasts, and events; generate notes and summaries while conversations are still in progress; build voice agents that need to understand users continuously; and create faster follow-up workflows for customer support, healthcare, sales, recruiting, and other high-volume spoken interactions.

Safety

safeguards and mitigations to help prevent misuse. We employ active classifiers over Realtime API sessions, meaning certain conversations can be halted if they are detected as violating our harmful content guidelines. Developers can also easily add their own additional safety guardrails using the Agents SDK.

Our usage policies prohibit repurposing or distributing outputs from our services for spam, deception, or other harmful purposes. Developers must also make it clear to end users when they're interacting with AI, unless it's already obvious from the context.

The Realtime API fully supports EU Data Residency for EU-based applications and is covered by our enterprise privacy commitments.

Pricing & availability

GPT-Realtime-2, GPT-Realtime-Translate and GPT-Realtime-Whisper are available in the Realtime API. GPT-Realtime-2 is priced at $32 / 1M audio input tokens ($0.40 for cached input tokens) and $64 / 1M audio output tokens. GPT-Realtime-Translate is priced at $0.034 per minute. GPT-Realtime-Whisper is priced at $0.017 per minute.

Get started

You can test the new realtime voice models in the Playground.

To start building, open this prompt in Codex to add GPT-Realtime-2 to an existing app or start a new

OpenAI

2026 API Platform

Author

OpenAI

Keep reading

View all

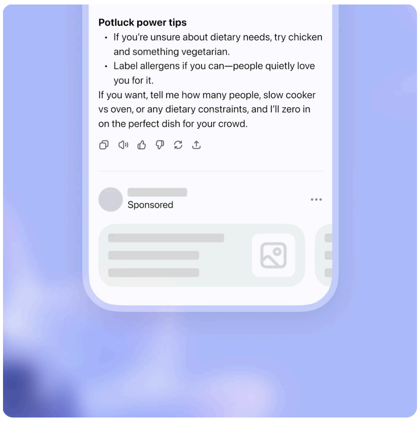

Testing ads in ChatGPT

Company May 7, 2026

GPT-5.5 Instant: smarter, clearer, and more personalized

Product May 5, 2026

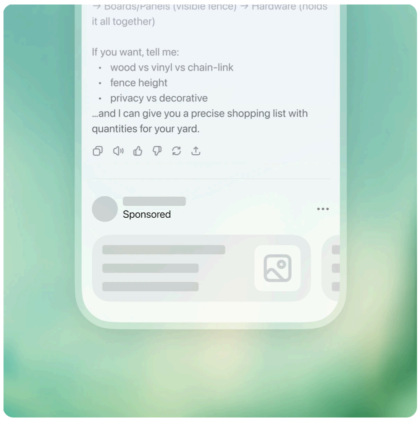

New ways to buy ChatGPT ads

Product May 5, 2026