Доверие — вот настоящее узкое место в агентном искусственном ...

Anthropic представляет Mythos для широкой аудитории с помощью Claude Fable 5 — самой мощной общедоступной модели.

Carl Franzen

10:19 по тихоокеанскому времени, 9 июня 2026 года

Источник: VentureBeat, создано с помощью ChatGPT-Images-2.0

Доверие — вот настоящее узкое место в агентном искусственном ...

Сегодня компания Anthropic выпустила две новые модели искусственного интеллекта — Claude Fable 5 и Claude Mythos 5 — что стало первым масштабным релизом мощных возможностей искусственного интеллекта «класса Mythos», которые ранее были доступны только организациям, участвующим в программе кибербезопасности Project Glasswing, о которой компания объявила два месяца назад.

По словам представителей компании, Fable 5, версия, которую большинство пользователей и разработчиков получат уже сегодня, превосходит все предыдущие общедоступные модели Claude по производительности в таких областях, как разработка программного обеспечения, интеллектуальная работа, визуализация, научные исследования и выполнение длительных задач.

Доверие — вот настоящее узкое место в агентном искусственном ...

Она превосходит существующие аналоги и лидирует почти по всем показателям, хотя предыдущая версия модели Claude Mythos Preview по-прежнему занимает первые места по использованию вычислительных ресурсов и междисциплинарному мышлению (см. сравнительную таблицу ниже и здесь).

Agentic coding FrontierCode (Diamond)	**29.3%** xhigh	—	13.4% xhigh	5.7% xhigh	—
Knowledge work GDPval-AA	**1932**	—	1890	1769	1314
Knowledge work vision GDP.pdf	**29.8%** no tools	—	22.5% no tools	24.9% no tools	16.7% no tools
Spatial reasoning Blueprint-Bench 2	**38.6%**	—	14.5%	36.2%	26.5%
Tool use AutomationBench	**17.4%**	—	15.5%	12.9%	9.6%
Computer use OSWorld-Verified	85.0%	**85.4%**	83.4%	78.7%	76.2%
Legal Legal Agent Benchmark	**13.3%**	—	10.4%	2.1%	0.0%
Multidisciplinary reasoning Humanity's Last Exam	**59.0%*** no tools	56.8% no tools	49.8% no tools	41.4% no tools	44.4% no tools
	64.5%* with tools	**64.7%** with tools	57.9% with tools	52.2% with tools	51.4% with tools
Biology BioMysteryBench	**46.1%*** hard	29.6% hard	40.0% hard	—	—
	83.9%* human solved	82.6% human solved	80.4% human solved	—	—
Agentic coding Terminal-Bench 2.1	**88.0%***	—	82.7%	83.4% Codex CLI	70.7% Gemini CLI
Cybersecurity ExploitBench (Cap%)	**78.0%***	69.0%	40.0%	34.0%	—
Health HealthBench Professional	**66.0%***	64.7%	56.9%	51.8%	—

Methodology: Reported scores are within a 1-3 percentage point difference for Claude Mythos 5 and Claude Fable 5. This table shows the higher score of the two. Starred (*) benchmarks show a larger difference due to our blocking safeguards for cybersecurity and biology-related questions. For these benchmarks, Claude Fable 5 performs closer to Claude Opus 4.8 due to fallbacks. See the system card for details.

Сравнительная таблица характеристик Claude Fable 5 и Mythos 5. Источник: Anthropic

Новый Claude Mythos 5, напротив, обладает более широкими возможностями, но и более ограниченными условиями доступности. Это обновленная версия предыдущей модели Mythos Preview, которая обладала схожими возможностями, но была доступна только в ограниченном количестве. Таким образом, в ней сняты некоторые ограничения, но официально она доступна только пользователям, одобренным Anthropic, в том числе партнерам Anthropic по кибербезопасности в рамках проекта Project Glasswing и некоторым исследователям в области биологии.

Доверие — вот настоящее узкое место в агентном искусственном ...

Ключевое отличие заключается в том, что универсальная модель Fable 5, основанная на тех же возможностях, что и Mythos, дополнена новыми средствами защиты. По словам представителей Anthropic, запросы, связанные с областями повышенного риска, включая кибербезопасность, биологию и химию, а также дистилляцию моделей, автоматически перенаправляются на Claude Opus 4.8, ранее бывшую флагманской универсальной моделью Anthropic. При этом пользователи получают уведомление. В случае с Mythos 5 такого не происходит.

По словам представителей компании, более 95 % сессий в Fable 5 полностью основаны на собственных ответах системы, без каких-либо резервных вариантов. После более чем 1000 часов тестирования внутренние и внешние специалисты по «красному» тестированию не обнаружили «универсальных уязвимостей».

Компания Anthropic сообщает, что Fable 5 уже доступна широкой публике на сайте, в приложениях и API, но Mythos 5 изначально будет доступна только тем пользователям, у которых уже есть доступ к более ранней версии Claude Mythos Preview.

Доверие — вот настоящее узкое место в агентном искусственном ...

Description	Anthropic's most capable widely released model, for the most demanding reasoning and long-horizon agentic work	Available through Project Glasswing. Successor to Claude Mythos Preview.
Claude API ID	`claude-fable-5`	`claude-mythos-5`
AWS Bedrock ID	`anthropic.claude-fable-5`	Limited availability
Vertex AI ID	`claude-fable-5`	Limited availability
Extended thinking	No	No
Adaptive thinking	Yes (always on)	Yes (always on)
Context window	1M tokens	1M tokens
Max output	128k tokens	128k tokens
Pricing	$10 / $50 per MTok (input / output)	$10 / $50 per MTok (input / output)

Claude Fable 5 is generally available on the Claude API, Claude Platform on AWS, Amazon Bedrock, Vertex AI, and Microsoft Foundry beginning June 9, 2026. Claude Mythos 5 is not generally available: it is offered in limited availability to approved customers in Project Glasswing, beginning the same day. For access, contact your Anthropic, AWS, or Google Cloud account team.

Скриншот с информацией об API Claude Fable 5 и Mythos 5 на сайте Anthropic.

Цены, доступ и непростое внедрение

Anthropic устанавливает цену на Fable 5 и Mythos 5 в размере 10 долларов за миллион входных токенов и 50 долларов за миллион выходных токенов. По словам представителей компании, это меньше половины стоимости Claude Mythos Preview, но все равно является самой высокой ценой среди основных моделей искусственного интеллекта, доступных по всему миру.

Обзор цен на API-интерфейсы моделей искусственного интеллекта от VentureBeat

deepseek-v4-flash	$0.14	$0.28	$0.42	ГлубОкий взгляд
deepseek-v4-pro	$0.435	$0.87	$1.305	ГлубОкий взгляд
Минимакс-M3	$0.30	$1.20	$1.50	Минимаксный
Gemini 3.1 Flash-Lite	$0.25	$1.50	$1.75	Google
Qwen3.7-Plus	$0.40	$1.60	$2.00	Облако Alibaba
MiMo-V2.5	$0.40	$2.00	$2.40	Xiaomi MiMo
Grok 4.3 (низкий контекст)	$1.25	$2.50	$3.75	xAI
ГЛМ-5	$1.00	$3.20	$4.20	Z.ai
Кими-К2.6	$0.95	$4.00	$4.95	Выстрел в Луну/Кими
GLM-5.1	$1.40	$4.40	$5.80	Z.ai
Grok 4.3 (высокий контекст)	$2.50	$5.00	$7.50	xAI
Qwen3.7-Max	$2.50	$7.50	$10.00	Облако Alibaba
Gemini 3.5 Flash	$1.50	$9.00	$10.50	Google
Предварительный просмотр Gemini 3.1 Pro (≤200 КБ)	$2.00	$12.00	$14.00	Google
GPT-5.4	$2.50	$15.00	$17.50	Открытый доступ
Предварительный просмотр Gemini 3.1 Pro (> 200 КБ)	$4.00	$18.00	$22.00	Google

Клод Фабль 5 / Клод Мифос 5	**$10.00**	**$50.00**	**$60.00**	**Антропный**

Для разработчиков Fable 5 доступна через Claude API как `claude-fable-5`. По словам представителей Anthropic, Fable 5 уже полностью доступна через Claude API и в корпоративных тарифных планах с оплатой по факту использования.

Для пользователей с подпиской процесс внедрения будет более сложным. По словам представителей Anthropic, с сегодняшнего дня и до 22 июня Fable 5 будет включена в тарифные планы Pro, Max, Team и Enterprise без дополнительной платы.

23 июня компания планирует исключить Fable 5 из этих тарифных планов, после чего для ее использования потребуются специальные кредиты. Anthropic заявляет, что намерена как можно скорее вернуть Fable 5 в стандартную комплектацию подписок.

Разница между Fable 5 и Mythos 5

Anthropic не позиционирует Fable 5 и Mythos 5 как две отдельные модели в привычном смысле «маленькая против большой». Скорее, они имеют одинаковый базовый уровень возможностей. Разница заключается в доступе *контроле* — то есть в том, насколько легко пользователям будет получить доступ к моделям, а также в встроенных ограничениях.

Как уже упоминалось, в Fable 5 появился новый уровень защиты, который выявляет определенные запросы с высоким уровнем риска, в том числе связанные с кибербезопасностью, биологией и химией, а также попытки перенести возможности модели в другие системы, и перенаправляет такие запросы в Claude Opus 4.8.

Mythos 5 снимает некоторые из этих ограничений для доверенных пользователей, работающих в одобренных доменах.

С практической точки зрения Mythos 5 более эффективен для работы с конфиденциальными данными в области кибернетики и биологии, поскольку он может справляться с задачами, с которыми не справляется Fable 5.

Доверие — вот настоящее узкое место в агентном искусственном ...

передовые модели с опасными возможностями двойного назначения: не предоставляя доступ ко всем возможностям всем желающим и не просто отклоняя рискованные запросы, а перенаправляя некоторые запросы в менее мощную модель, оставляя более мощную модель доступной для выполнения большинства повседневных задач.

Значительное улучшение в области автономного программирования

Для корпоративных клиентов наиболее очевидным вариантом использования, скорее всего, станет разработка программного обеспечения. По словам представителей Anthropic, Fable 5 может работать без участия человека дольше и с большей автономностью, чем предыдущие модели Claude. Именно такие возможности нужны компаниям, если они хотят, чтобы ИИ-агенты не только дополняли код или отвечали на вопросы разработчиков.

По данным Anthropic, на **тестировании SWE-bench Pro, которое оценивает способность модели выполнять сложные задачи в области разработки программного обеспечения, Fable 5 и Mythos 5 набрали 80,3%**, значительно опередив новейшую и самую совершенную универсальную модель GPT-5.5 от OpenAI, которая набрала 58,6%.

Согласно таблице результатов, приведенной в материалах Anthropic, в тесте FrontierCode Diamond от Cognition, который оценивает качество и удобство сопровождения агентного кода, модели набрали 29,3 % баллов по сравнению с 13,4 % у Claude Opus 4.8 и 5,7 % у GPT-5.5.

По данным Anthropic, Fable 5 показала самые высокие результаты среди передовых моделей на FrontierCode даже при средних вычислительных затратах. Это позволяет предположить, что модель может выдавать более качественные результаты при кодировании, не требуя максимальных вычислительных мощностей.

Самый впечатляющий пример — от компании Stripe. По словам представителей Anthropic, Stripe протестировала Fable 5 на базе кода Ruby, состоящей из 50 миллионов строк, и обнаружила, что модель выполнила миграцию всей баз кода за один день, в то время как вручную на это у команды ушло бы больш двух месяцев. В Stripe сказали: «Fable 5 сокращает месяцы разработки до a class of long-horizon problems that were out of reach for earlier models." Replit said Fable 5 is the highest-performing model it has tested on ViBench, its end-to-end "vibe-coding" benchmark, and that it builds apps in less time with fewer tokens. Figma said Fable 5 is "a clear step forward on agentic coding and prototyping."

This is the enterprise shift Anthropic is trying to sell: AI coding systems that can take on larger units of work, not just individual tickets. That could include codebase migrations, app prototyping, pull request review, test generation, debugging across unfamiliar tools, user interface design and multi-step internal software projects.

Base44 said, "Fable 5 is much deeper and better at one-shotting full apps, and its tool calling is excellent." Genspark said, "Fable 5 came out #1 on our evals, winning head-to-head against every model we tested. It was significantly stronger on the hardest tasks in the set — UI design and game coding." Rakuten said, "At the highest effort, Fable 5 reflects on and validates its own work. For us, that's what makes highly autonomous operations possible — the extra thinking pays for itself."

For CTOs and engineering leaders, that suggests the model's value may come less from raw code generation and more from sustained execution: understanding an intent, planning steps, calling tools, checking its own work and continuing through a task without constant human steering.

Knowledge work, finance, legal and operations

Anthropic is also positioning Fable 5 as a stronger model for enterprise knowledge work. On GDPval-AA, Anthropic reports a score of 1932 for Fable 5 and Mythos 5, compared with 1890 for Claude Opus 4.8, 1769 for GPT-5.5 and 1314 for Gemini 3.1 Pro.

On GDPpdf, a benchmark focused on visual document reasoning, Fable 5 and Mythos 5 score 29.8% without tools, compared with 22.5% for Opus 4.8, 24.9% for GPT-5.5 and 16.7% for Gemini 3.1 Pro.

That matters for enterprises because much of corporate work still lives in messy documents: PDFs, spreadsheets, charts, reports, contracts, filings, slide decks and screenshots. Anthropic says Fable 5 shows gains in document-based reasoning, chart and table interpretation and complex problem solving.

solving.

The finance examples are notable because they point to AI agents moving beyond summarization into higher-stakes analytical workflows.

IMC said Fable 5 "aced our trading-analysis evaluations nearly across the board: factual lookup, conceptual reasoning, root-cause analysis, expected-value analysis." Optiver said the model was stronger than Opus 4.8 on its trading benchmark and "remarkably consistent," scoring identically across repeated runs. Balyasny Asset Management said Fable 5 was the strongest finance-first model it had tested.

Legal and operations teams may also see immediate impact. Crosby Legal said, "Fable 5 feels materially different. In blind review, our lawyers found its redlines matched or beat our current model every time." Notion said the model can take work "you'd chip away at all afternoon" and turn messy notes into a functioning project plan. Zapier said Fable 5 is the new leader on AutomationBench and is more autonomous than Opus 4.8: "Where Opus stops to ask, Fable 5 keeps looking."

For enterprise software vendors, that points toward more capable embedded agents in workflow products: agents that can review a contract, update a project plan, assemble a spreadsheet, inspect a chart, file a ticket, run a query, call an internal API and keep going until the work is complete.

Vision and interface understanding

Anthropic says Fable 5 is also its strongest vision model. In its launch materials, the company says the model can extract precise numbers from detailed scientific figures and complete vision-based tasks such as rebuilding a web app's source code from screenshots alone.

That has immediate implications for enterprise automation. Many business processes still depend on visual interfaces that are not cleanly exposed through APIs: dashboards, PDFs, forms, legacy apps, screenshots, scans and image-heavy reports. A stronger vision model could help agents operate across those environments with less custom integration work.

Anthropic also says Fable 5 needs less scaffolding than previous Claude models. A example, the company says earlier Claude models struggled to play Pokémon Fire even with extra tools, while Fable 5 impressively beat the game using a minimal vision-

The point is not gaming itself, but the broader agentic skill: reading a visual environment, remembering progress, deciding what to do next and executing over a long horizon.

In another internal test, Anthropic says it had the model play the deck-building game Slay the Spire with access to persistent file-based memory. The company says persistent memory improved Fable 5's performance three times more than it improved Opus 4.8's, and that Fable reached the game's final act three times more often. For enterprise users, this suggests Fable 5 may make better use of notes, logs and stored context during multi-step work.

That could matter for internal agents that operate over days or weeks: sales operations agents that track account research, engineering agents that manage migrations, finance agents that update models, or support agents that remember what they tried across many turns.

From restricted cyber model to general-purpose enterprise AI

The announcement follows Anthropic's April 2025 rollout of Claude Mythos Preview through Project Glasswing, a restricted program for cyber defenders, critical infrastructure providers and major software maintainers. Anthropic created Glasswing after internal evaluations showed Mythos-class models could find and exploit software vulnerabilities at a level that raised meaningful misuse concerns.

Доверие — вот настоящее узкое место в агентном искусственном ...

banking officials and regulators in the U.S. and U.K. because of fears that AI-accelerated cyberattacks could threaten payment systems and broader financial stability.

The reaction has not been limited to alarm. Governments also want access: Reuters reported that South Korea's national internet security agency had secured Mythos access through Project Glasswing, reflecting a broader geopolitical race to use frontier AI for national cyber defense. At the same time, Anthropic has faced scrutiny over whether it can safely gate the very capabilities it says are too risky for general release. The Verge reported that unauthorized users accessed Mythos after its limited rollout, calling the incident damaging for a company that has built its brand around responsible AI.

Critics have also questioned whether Anthropic's warning-heavy framing risks becoming a form of market positioning, since it casts the company as both the source of the new capability and the gatekeeper deciding which governments, companies and researchers get to use it.

With Fable 5, Anthropic is leaning into its gatekeeper role, attempting to separate the general enterprise value of a Mythos-class model from the riskiest parts of its capability profile. The company says Fable 5 can handle software engineering, research, visual reasoning, document analysis and long-running agentic workflows, while classifiers block or reroute requests that could provide what Anthropic calls "uplift" to malicious actors.

Those classifiers cover three main areas.

1. Cybersecurity, where Anthropic says Mythos-class models can discover and exploit vulnerabilities and perform broader "agentic hacking" tasks such as reconnaissance, discovery and lateral movement.

2. Biology and chemistry, where the company says the same reasoning that can help researchers design therapies could also help well-resourced malicious actors pursue dangerous biological work.

3. Model distillation, where Anthropic says users may try to extract Claude's capabilities to train competing models, including models that could be released without similar safeguards.

When Fable 5's classifiers detect one of those categories, the response is automati handled by Claude Opus 4.8. Anthropic says users will be told when this happens.

company says an external bug bounty produced no universal jailbreaks after more than 1,000 hours of testing, and external red-teaming organizations also failed to find a universal jailbreak. One external partner found that Fable 5 complied with zero harmful single-turn cyber requests related to planning cyberattacks, exploit development or defense evasion, even when prompts used any of 30 public jailbreak techniques, according to Anthropic.

The company is still acknowledging tradeoffs. Anthropic says the safeguards are deliberately cautious and may sometimes trigger on benign requests. That could frustrate security professionals, biology researchers and advanced enterprise users whose legitimate work overlaps with the blocked categories. The company says it plans to reduce false positives over time.

Mythos 5 and the restricted frontier

While Fable 5 is the broad commercial launch, Mythos 5 is the model to watch for enterprises operating in security, critical infrastructure and life sciences.

The company says all users with Claude Mythos Preview access can upgrade to Mythos 5 beginning today. It plans to expand access through a trusted access program, in collaboration with the U.S. government.

The distinction is important for sectors where the blocked capabilities are not edge cases but core workflows. A security team may need to reproduce vulnerabilities, test exploitability, analyze lateral movement or simulate attacker behavior in a controlled environment. A biology research team may need to reason through molecular design workflows that would trigger general-use safeguards. Fable 5 is not designed to give every user unrestricted access to those capabilities; Mythos 5 is designed for vetted users who need them.

Anthropic says Mythos 5 has the strongest cybersecurity capabilities of any model in the world. In the company's benchmark table, the model family scores 78.0% on ExploitBench, compared with 69.0% for Claude Mythos Preview, 40.0% for Opus 4.8 and 34.0% for GPT-5.5. On CyberGym, Anthropic's chart shows Mythos 5 at 83.8%, slightly ahead of Mythos Preview at 83.1% and far above Opus 4.8 with default safeguards.

The company is making a similar argument in biology. Anthropic says Mythos-cl
models outperform dedicated protein language models on a task involving adeno-

the drug design process by about tenfold. In one example, the company says Mythos 5, using protein design and bioinformatics tools without human assistance, matched or beat skilled human operators by choosing binding sites, selecting and running tools, and recovering from failures. Anthropic says nine of 14 protein targets in the study produced strong candidates for drug design that it is now investigating.

The company also says Mythos 5 produced novel molecular biology hypotheses that Anthropic scientists preferred over Opus-class model hypotheses about 80% of the time in blinded comparisons. Anthropic says several of those ideas have advanced to experimental evaluation, and one hypothesis involving an E. coli protein was later corroborated by an independent lab working on the same problem.

Those claims are potentially significant, but they should be treated carefully until more details are published. Anthropic says it intends to publish additional results in the coming months. For now, the strongest enterprise implication is directional: the company believes its highest-end models can already perform parts of scientific research workflows with less human intervention than prior systems.

New, longer data retention requirement

The company also introduced a new data-retention policy for Mythos-class models. Anthropic says it will require 30-day retention for all traffic on Fable 5, Mythos 5 and future models with similar or higher capability levels, across both first-party and third-party surfaces. The company says it will not use that data to train new Claude models or for non-safety purposes, and says it has added privacy protections including logging human access and deleting the data after 30 days in almost all cases.

That policy may become one of the most important enterprise buying questions around Fable 5. Many businesses want frontier AI capability but also want strict control over data retention, especially in regulated sectors. Anthropic's position is that stronger monitoring is necessary for models with this level of capability. Enterprise customers will have to decide whether the capability gain justifies the retention requirement.

Enterprise implications

The broader enterprise significance of Fable 5 is that Anthropic is trying to commercialize a more autonomous class of AI model without exposing all of its larger tasks: code migrations, refactors, UI builds, test writing, bug fixing, documentation, internal tooling and multi-step app creation.

For knowledge-work-heavy enterprises, Fable 5 could make AI more useful in workflows where earlier models were too brittle: finance research, spreadsheet analysis, legal redlines, procurement review, board materials, market research, sales operations and project planning. The main gain is not just better answers; it is fewer turns, fewer corrections and more ability to keep working through ambiguity.

For security teams, the launch is more complicated. Most organizations will get Fable 5, not unrestricted Mythos 5. That means they may see stronger general coding and analysis, but not full access to the cyber capabilities Anthropic considers risky. Trusted defenders inside Project Glasswing will get Mythos 5, giving them a more direct way to use the model for vulnerability discovery and defensive testing.

For life sciences companies, the pattern is similar. Fable 5 may help with general research, literature analysis, data interpretation and scientific reasoning, but the more sensitive biological capabilities will be restricted. Anthropic is effectively creating a separate access path for vetted researchers whose work requires capabilities that could be dangerous in the wrong hands.

The launch also raises competitive pressure across the AI industry. Anthropic is claiming state-of-the-art results across agentic coding, knowledge work, vision, cybersecurity, legal reasoning, spatial reasoning and health benchmarks. But the more strategically important claim may be that it has found a workable release mechanism for models above its Opus class. If Fable 5's safeguards hold up under real-world use, Anthropic will argue it can bring more powerful models to market sooner without fully opening the riskiest capabilities.

That is still a large "if." The enterprise market will test not only Fable 5's benchmark performance, but also its reliability, false-positive rate, data-retention tradeoffs and cost at scale. A model that can complete more work autonomously can also burn more tokens, trigger more governance questions and create new review burdens for teams that must verify its output.

Still, today's launch marks a clear shift in the Claude lineup. Opus is no longer Anthropic's top commercial capability tier. Mythos-class models now sit above it. Fable 5 is the first version of that tier for general users; Mythos 5 is the restricted version for trusted high-risk work. Together, they show how Anthropic plans to push frontie

Доверие — вот настоящее узкое место в агентном искусственном ...

Subscribe to get latest news!

Deep insights for enterprise AI, data, and security leaders

VB Daily
AI Weekly
AGI Weekly
Security Weekly
Data Infrastructure Weekly
VB Events
All of them

Доверие — вот настоящее узкое место в агентном искусственном ...

Get updates

More

Доверие — вот настоящее узкое место в агентном искусственном …

Satya Nadella warns that AI could hollow out entire industries, echoing the damage done by globalization

Microsoft CEO Satya Nadella published a sweeping essay on Sunday laying out what he describes as the defining economic challenge of the AI era: the risk that a handful of frontier models will absorb the expertise of entire industries and commoditize it,...

Michael Nuñez

June 15, 2026

Доверие — вот настоящее узкое место в агентном искусственном ...

model training or fine-tuning unless the client provides explicit opt-in consent.

Carl Franzen

June 15, 2026

Anthropic blocks all public access to Claude Fable 5, Mythos 5 following US government order — what enterprises should do

Enterprises can no longer afford — from an operational reliability standpoint — to run critical workflows on any single AI model or even provider.

Carl Franzen

June 13, 2026

Доверие — вот настоящее узкое место в агентном искусственном ...

Kimi K2.7-Code cuts thinking tokens 30% — but practitioners say the benchmarks don't check out

K2.7-Code authors code directly instead of wrapping libraries — more honest, but two kernels failed and the MoE result regressed from K2.6.

Sean Michael Kerner

June 13, 2026

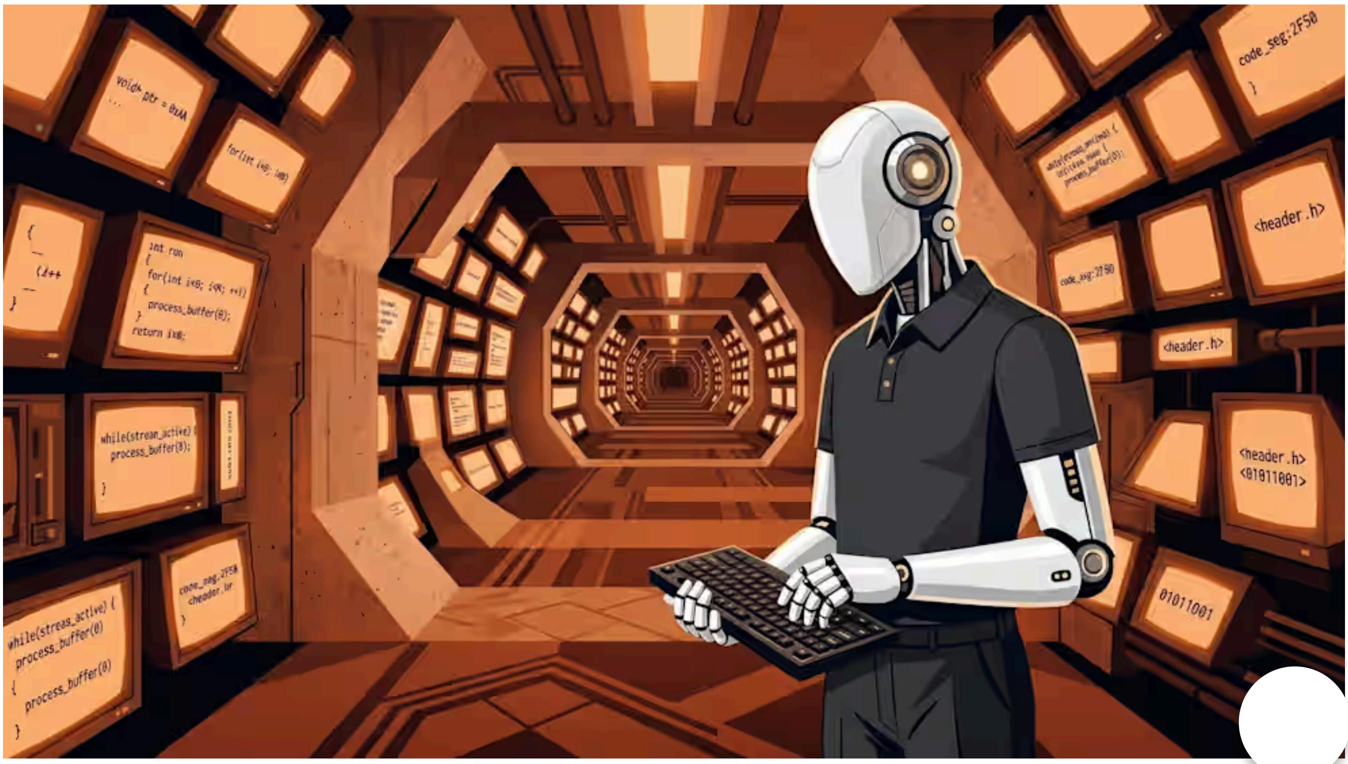

Доверие — вот настоящее узкое место в агентном искусственном …

Carl Franzen

June 12, 2026

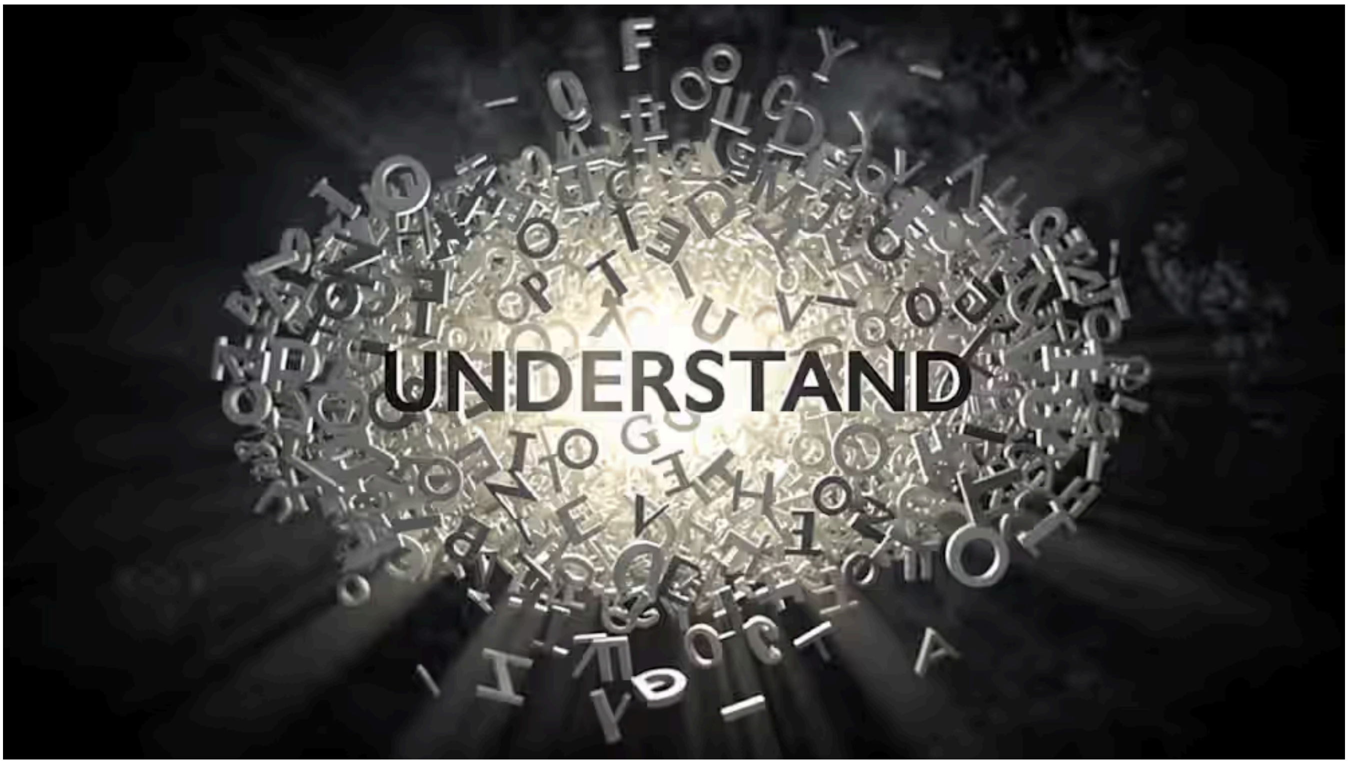

Google's DiffusionGemma generates 256 tokens in parallel and self-corrects as it goes

It writes a full block of tokens at once instead of one at a time, and fixes its own errors mid-generation. Fast on consumer GPUs, weaker on open-ended tasks.

Sean Michael Kerner

June 11, 2026

Доверие — вот настоящее узкое место в агентном искусственном ...

Surprise upset: GPT-5.5 beats Claude Fable 5 on brutal new Agents' Last Exam benchmark

The victory of GPT-5.5 aligns with recent third-party analysis suggesting that OpenAI's models are currently superior at strictly adhering to multi-part, complex prompts

Carl Franzen

June 11, 2026

Доверие — вот настоящее узкое место в агентном искусственном ...

Ben Dickson

June 11, 2026

Anthropic CEO calls for FAA-style regulation of powerful AI models: what enterprises should know

To prepare for this shift, enterprises must first decouple their AI strategies from single-vendor dependencies. If a flagship model is suddenly blocked or recalled under the proposed FAA-style regulatory powers

Carl Franzen

June 10, 2026

Доверие — вот настоящее узкое место в агентном искусственном ...

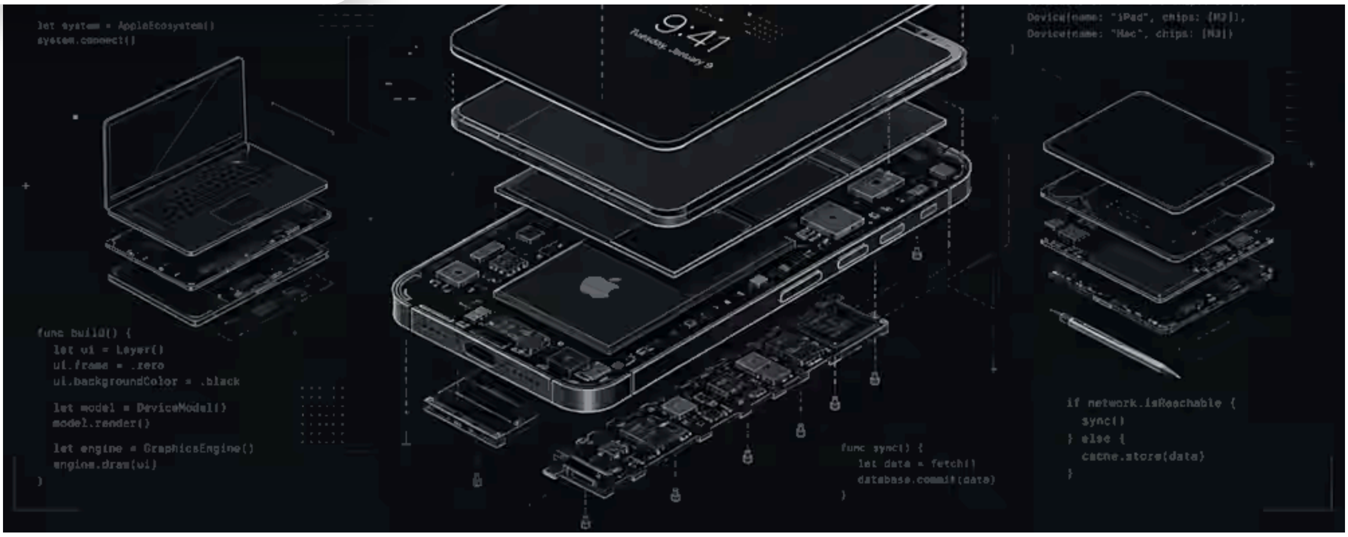

Apple’s new Siri AI is more than just a smarter assistant — it's a new enterprise app layer

For enterprise technology leaders, it means Apple’s devices could soon include a native AI assistant that can act across business workflows

Carl Franzen

June 10, 2026

Доверие — вот настоящее узкое место в агентном искусственном …

Sean Michael Kerner

June 10, 2026

On-device AI agents hit a hard memory limit. Apple's new architecture routes around it.

Apple's flash-routing architecture puts 20B parameters on-device without touching DRAM. For enterprises locked out of cloud inference, that's a significant new option.

Sean Michael Kerner

June 9, 2026

Доверие — вот настоящее узкое место в агентном искусственном ...

Press Releases

Contact Us

Advertise

Share a News Tip

Contribute

Privacy Policy

Terms of Service

Consent Preferences

Do Not Sell or Share My Personal Information

Limit the Use Of My Sensitive Personal Information