

← Вернуться к блогам

Новости искусственного интеллекта

Новости искусственного интеллекта 11 июня 2026 года: 12 главных событий дня

📅 11 июня 2026 года ⌚ 19 минут чтения



Поделиться:



Новости об искусственном интеллекте от 11 июня 2026 года: 12 главных событий, о которых нужно знать сегодня

Индустрия искусственного интеллекта не дремлет, и 11 июня 2026 года — тому подтверждение. Только сегодня компания SpaceX

провела крупнейшее в истории IPO по цене 135 долларов за акцию на бирже Nasdaq, компания OpenAI завершает сделку по интеграции своих моделей в Oracle Cloud, а гонка между Anthropic, OpenAI и Google как никогда напряжённая. Если вы не спали последние 48 часов, то вот вам все важные новости, с источниками и проверками.

[Центр новостей и трендов в сфере искусственного интеллекта](#) — это самый быстрый способ быть в курсе всех важных событий в OpenAI, Anthropic, Google и экосистеме с открытым исходным кодом.

1. Цена IPO SpaceX 11 июня — 135 долларов за акцию. Крупнейшее IPO в истории

Сегодня, 11 июня 2026 года, после закрытия торгов на бирже SpaceX устанавливает фиксированную цену в 135 долларов за акцию, что позволит привлечь около 75 миллиардов долларов и оценить компанию примерно в 1,77 триллиона долларов. Если все пройдет по плану, это будет крупнейшее первичное публичное размещение акций в истории — оно превзойдет листинг Saudi Aramco в 2019 году, когда компания была оценена в 35,4 миллиарда долларов при рыночной капитализации в 320 миллиардов долларов.

Торги начнутся на бирже Nasdaq под тикером SPCX в пятницу, 12 июня. Goldman Sachs возглавляет синдикат из 21 банка, при этом 30% акций будут распределены среди розничных инвесторов через Robinhood, Fidelity и Charles Schwab. Доля розничных инвесторов примерно в три раза превышает стандартные 5–10%.

Илон Маск не продаст ни одной акции. Двухклассовая структура компании гарантирует, что он сохранит полный контроль над голосованием. Источники дохода компании: только Starlink в 2025 году принёс 11,4 млрд долларов выручки и 4,4 млрд долларов операционной прибыли. В феврале 2026 года SpaceX поглотила

компанию xAI, выкупив все её акции, и оценила объединённую компанию в 1,25 триллиона долларов.

Горячий анализ: при 100-кратном превышении прогнозируемого дохода цена SPCX рассчитана на десять лет работы без права на ошибку. В первый день торгов акции компаний, ориентированных на розничную торговлю, обычно демонстрируют резкие колебания. Те, кто покупает акции в первый день торгов, инвестируют не в фундаментальные показатели SpaceX, а делают ставку на краткосрочный импульс.

2. OpenAI + Oracle: доступ к моделям искусственного интеллекта теперь можно получить с помощью облачных кредитов Oracle

11 июня 2026 года компания OpenAI официально объявила, что корпоративные клиенты теперь могут получить доступ к передовым моделям искусственного интеллекта и Codex с помощью существующих универсальных кредитов Oracle (Universal Credits, UCM). [Объявление о партнерстве OpenAI и Oracle](#) позволяет клиентам Oracle Cloud Infrastructure (OCI) использовать технологии OpenAI без создания отдельного канала закупок.

This is a significant enterprise move. Companies with existing Oracle cloud commitments - often multi-year, pre-negotiated contracts running into the tens or hundreds of millions - can now route AI workloads to OpenAI models without a new vendor relationship. The availability is expected to go live in the coming weeks; Oracle customers should contact their sales representative for timing details.

Context: OpenAI and Oracle have been deeply connected since the Stargate project launched in January 2026. Stargate involves a \$500 billion infrastructure commitment with Oracle building data centers in Texas, New Mexico, Wisconsin, Michigan, and other states. This latest move turns that infrastructure investment into an enterprise distribution channel.

My read: This partnership does more for enterprise AI adoption than any benchmark improvement. Procurement friction is the real barrier to enterprise AI deployment, not model quality. Removing it is a bigger deal than it sounds.

3. Anthropic Files Confidential S-1 with SEC at \$965B Valuation

Anthropic confidentially submitted a draft registration statement on Form S-1 with the US Securities and Exchange Commission on June 1, 2026, targeting a potential public offering near its most recent \$965 billion post-money valuation. The company had just closed a \$65 billion Series H round - the largest private AI fundraise in history - days before the filing.

Goldman Sachs, Morgan Stanley, and JPMorgan are leading the IPO process. Anthropic stated: "This gives us the option to go public after the SEC completes its review." The company reported a run-rate revenue of \$47 billion as of May 2026 and expects its annualized run rate to surpass \$50 billion by end of June 2026. Anthropic is also approaching its first profitable quarter.

Wedbush Securities analyst Dan Ives called the move "a major step to get ahead of OpenAI" and described the cluster of AI IPOs (SpaceX, Anthropic, OpenAI) as "an opening of the floodgates for the IPO market." All three companies are working with investment banks that have advised on the importance of being first to list and set the terms for how investors categorize the AI sector.

For a complete breakdown of Anthropic's model lineup and Claude's growing competitive position, see the [Claude AI Complete Hub](#) covering Opus, Sonnet, Haiku, and Mythos.

4. OpenAI Confidentially Files Its Own S-1 at \$852B Valuation

One week after Anthropic's S-1, OpenAI announced on June 8, 2026, that it had also confidentially submitted a draft S-1 to the SEC. In a characteristically transparent move, the company posted its own announcement: "We recently submitted a confidential S-1. We expect it to leak so we're just announcing it."

OpenAI is working with Goldman Sachs and Morgan Stanley as lead underwriters. At its March 2026 \$122 billion funding round, the company was valued at \$852 billion with over \$20 billion in annual recurring revenue for 2025 - a tripling of figures year-over-year since 2023. Internal documents, however, project a \$14 billion loss in 2026 with profitability not expected until 2029.

CEO Sam Altman framed the filing as the beginning of a third phase for OpenAI: "The economy is beginning to reshape around AI." Analysts at CNBC and Enterprise DNA expect a public listing valuation exceeding \$1 trillion - roughly 4x Alibaba's 2014 debut valuation. A potential listing window of September to November 2026 has been cited by multiple sources, though no timeline is confirmed.

The filing came two days after a jury dismissed Elon Musk's lawsuit against OpenAI on statute of limitations grounds - a clean legal slate entering the public market process.

5. ChatGPT Crosses 1 Billion Monthly Active Users

ChatGPT became the fastest application in history to reach 1 billion global monthly active users, crossing the milestone in May 2026 - roughly three years after its November 2022 launch. Sensor Tower confirmed the figure, reported by Reuters on June 2, 2026. The achievement outpaced the growth trajectories of Google Maps, TikTok, Instagram, and YouTube.

The numbers behind the headline are equally striking. As of Q2 2026, ChatGPT had 900 million weekly active users, while Claude had 56 million global monthly active app users. However, Claude's year-over-year MAU growth of approximately 640% dramatically outpaces ChatGPT's 62% growth rate. Sensor Tower also found that US ChatGPT

users who installed the Claude app in Q1 2026 spent 5% less time on ChatGPT one month later - a small but measurable shift in attention.

One billion users is not a ceiling - it is a floor. The more interesting metric is Claude's 640% YoY growth. That trajectory, if sustained for even two more years, rewrites the competitive map entirely. ChatGPT won the awareness war. The usage war is still being fought.

6. Claude Opus 4.8 Launches with Dynamic Workflows and Honesty Gains

Anthropic launched Claude Opus 4.8 on May 28, 2026, describing it as "a modest but tangible improvement" over Opus 4.7. The [Anthropic official announcement](#) highlighted improvements across agentic coding, multidisciplinary reasoning, agentic computer use, and professional workflows.

The headline improvement is honesty. Opus 4.8 is approximately four times less likely than its predecessor to allow flaws in code it generates to pass unremarked. It is also more likely to flag uncertainty and less likely to make unsupported claims. These honesty metrics matter because agentic systems operating autonomously for extended periods need to be self-correcting, not confident-and-wrong.

Alongside the model, Anthropic released Dynamic Workflows (research preview) for Claude Code - a feature that lets the model plan work, run hundreds of parallel subagents within a single session, and verify outputs before returning results. It is available on Enterprise, Team, and Max plans. Fast mode for Opus 4.8 now runs at 2.5x the speed and is three times cheaper than fast mode for previous Opus models.

Anthropic also noted that Opus 4.8 performs at a level similar to Claude Mythos Preview on prosocial benchmarks - measures of whether the model acts in accordance with user interests and instructions. The company has signaled Mythos-class models will be available to all customers "in the coming weeks."

7. DeepSeek V4 Pro Makes 75% Price Cut Permanent

DeepSeek officially made its 75% price discount on V4 Pro permanent on May 23, 2026, after the promotional period was set to expire on May 31. The new permanent pricing for DeepSeek V4 Pro runs from \$0.003625 to \$0.87 per one million tokens - down from the previous range of \$0.0145 to \$3.48 per million tokens.

To put that in competitive context: OpenAI's GPT-5 charges \$2.50 per million input tokens and \$10 per million output tokens. Anthropic's Claude Opus 4.7 is priced at \$5 input and \$25 output. Google's Gemini 3.5 Flash sits at \$0.15 input and \$0.60 output. DeepSeek V4 Pro's permanent pricing sits below every major competitor, and it runs on a 1M-token context window.

DeepSeek's explanation for the permanence: "V4-Pro was engineered to cut the cost of long-context inference, reportedly running at roughly a quarter of the single-token compute and a tenth of the memory footprint of its predecessor at very long context. This is why the price cut is permanent rather than promotional. It is not a discount." The model runs on Huawei's Ascend 950 chips, reducing its exposure to Nvidia supply constraints and US export controls.

The era of high-margin AI tokens may be ending faster than anyone expected. DeepSeek is running the Amazon retail playbook: lose money on price, win the customer, build the moat. The structural question for Western AI labs is whether they can match Chinese pricing while funding \$100B+ infrastructure buildouts.

8. Google Gemini 3.5 Pro - Still Pending, Still Imminent

Google announced Gemini 3.5 Pro at Google I/O on May 19, 2026, with Sundar Pichai telling the audience: "Give us until next month to get it to you." That month is now June 2026. As of June 11, Gemini 3.5 Pro has not yet shipped to general availability, though it has been in limited Vertex AI preview and internal use.

What is confirmed about the model: a 2 million token context window, a "Deep Think" reasoning mode, and frontier multimodal understanding. In Google's lineup, 3.5 Pro absorbs the use cases previously routed to the Ultra tier - the hardest reasoning tasks, deep multimodal workloads, and very long context. Gemini 3.5 Flash, which did ship on May 19, already scores 76.2% on Terminal-Bench 2.1 and 83.6% on MCP Atlas, beating Gemini 3.1 Pro on both.

Prediction markets on Polymarket are pricing in a late-June 2026 launch with probabilities clustered across mid-to-late June. If Pro extends the Flash performance gap in the same proportions, [the competitive implications for Claude Opus and GPT-5.5 are significant](#). The benchmark race in H2 2026 looks genuinely competitive for the first time since GPT-4.

9. Microsoft Build 2026 - 7 MAI Models and a New Era of Independence

Microsoft Build 2026 opened on June 2 at Fort Mason Center in San Francisco with a single clear message: Microsoft is no longer dependent on OpenAI for AI capabilities. The company unveiled seven in-house MAI models trained entirely without OpenAI involvement - the first full exercise of rights restored in April 2026 when restrictions in the Microsoft-OpenAI partnership were lifted.

The seven MAI models announced: MAI-Thinking-1 (reasoning flagship), MAI-Code-1-Flash (coding model, 60% fewer tokens than comparable models), MAI-Image-2.5 (image generation, ranked above Gemini in Arena leaderboard), MAI-Transcribe-1.5 (state-of-the-art transcription across 43 languages, 5x faster than rivals), MAI-Voice-2 (speech generation), and two additional reasoning-specialized models for Azure AI Foundry.

MAI-Code-1-Flash is live in the GitHub Copilot model picker in VS Code as of June 2, available on all paid tiers. Under the usage-based billing model that started June 1, its 60% token efficiency advantage translates directly

to credit savings. Microsoft also announced Project Solara for agent devices, Agent 365, Work IQ APIs, and the Majorana 2 quantum chip.

The Build 2026 announcements also coincided with the single largest government contract in Microsoft history - an AI capabilities deal with the US Department of Defense covering 2.1 million service members and 770,000 civilian staff.

10. GitHub Copilot Token Billing Goes Live - Developer Backlash Erupts

On June 1, 2026, GitHub Copilot switched every plan - Free, Pro (\$10/month), Pro+ (\$39/month), Business (\$19/user/month), and Enterprise (\$39/user/month) - from flat-rate subscriptions to token-based billing denominated in GitHub AI Credits, where 1 credit equals \$0.01 USD.

The immediate reaction from developers was fierce. The core complaint: a single agentic coding session now consumes \$30 to \$40 in credits, three to four times a Pro subscriber's entire \$10 monthly allotment. On Reddit and Hacker News, developers accused Microsoft of actively building features that encourage high token consumption, then penalizing users for using them. One Reddit user wrote: "Microsoft provided this billing method and they kept making it easier to burn through massive numbers of tokens while spawning dozens or even hundreds of sub-agents."

Microsoft's counter-argument: flat-rate was unsustainable. Heavy users were being subsidized by lighter users. Token billing aligns pricing with actual usage. Business and Enterprise customers receive promotional included credits through September 1, 2026 - a grace period that gives teams time to build cost monitoring infrastructure.

MAI-Code-1-Flash's arrival at Build 2026 one day after the billing change was not coincidental. At 60% lower token consumption, it directly

addresses the credit burn problem for the users most affected. The all-you-can-eat era of AI coding assistance for individual developers is over.

11. Anthropic Raises \$65B Series H - Surpasses OpenAI in Private Valuation

On May 28, 2026, Bloomberg confirmed that Anthropic closed a \$65 billion Series H funding round at a \$965 billion post-money valuation - surpassing OpenAI's \$852 billion private market valuation for the first time and making Anthropic the most valuable private AI company in the world. The round was co-led by Sequoia, Dragoneer, Altimeter, and Greenoaks, with additional equity from Microsoft, NVIDIA, Founders Fund, and General Catalyst.

Alongside the equity round, Apollo Global Management and Blackstone structured a \$36 billion private credit deal to purchase Google custom TPU chips on behalf of Anthropic. The chips will be leased to Anthropic for use at data centers in New York, Texas, Louisiana, and Indiana. This infrastructure commitment is central to Anthropic's ability to scale Claude deployment while preparing for a public listing.

Anthropic's run-rate revenue crossed \$47 billion in May 2026, and the company expects annualized run rate revenue to surpass \$50 billion by end of June. It is approaching its first profitable quarter - a milestone that transforms the IPO narrative from "ambitious AI lab" to "high-growth, near-profitable enterprise platform."

12. AI Model Price War Intensifies Across the Entire Industry

June 2026 is when the AI price war went from competitive to structural. DeepSeek V4 Pro's permanent 75% price cut sits below every Western frontier model. Gemini 3.5 Flash at \$1.50/\$9 per million tokens is positioned as the "default tier for agentic-first developers" - a deliberate repositioning away from cheap alternative and toward primary workhorse. OpenAI's pivot toward consumer platform features and

advertising signals that API token revenue alone may not sustain its \$852 billion valuation.

The pressure point is not the top-tier enterprise contracts. Those still go to Claude and GPT-5 based on capability and compliance. The pressure is on the mid-market: startups, agencies, and independent developers who are price-sensitive and increasingly willing to route workloads to whatever is cheapest and good enough. DeepSeek is winning that segment aggressively.

For a head-to-head comparison of the current model landscape including pricing and benchmark performance, see the [Best AI Models and Leaderboards hub](#) for the latest rankings.

The honest contrarian take: commoditization of inference is actually good for the AI ecosystem. Lower token costs mean more developers can build more applications, which generates more demand, which funds the next frontier model generation. The labs complaining about margin compression are the same ones raising tens of billions in capital. The price war is a feature, not a bug.

Frequently Asked Questions

What is the biggest AI news on June 11, 2026?

The most significant development on June 11, 2026 is SpaceX pricing its IPO at \$135 per share after market close, targeting a \$75 billion raise and \$1.77 trillion valuation - which would be the largest IPO in recorded history. Simultaneously, OpenAI announced that its frontier models and Codex are now accessible through Oracle Cloud Infrastructure using existing Oracle Universal Credits.

Has OpenAI filed for an IPO?

Yes. OpenAI confidentially submitted a draft S-1 registration statement to the US Securities and Exchange Commission on June 8, 2026 (announced

publicly the same day). Goldman Sachs and Morgan Stanley are leading the process. The company has not set a timeline, noting that going public "may be a while," but that the filing gives it the option to list sooner if market conditions are favorable. Analysts point to a September to November 2026 window as most likely.

What is Claude Opus 4.8 and what is new?

Claude Opus 4.8 is Anthropic's latest flagship model, launched May 28, 2026 at the same price as Opus 4.7. Key improvements include four times lower likelihood of allowing code flaws to pass unremarked, stronger honesty and uncertainty flagging, and better agentic performance across coding, reasoning, and financial analysis. It launched alongside Dynamic Workflows for Claude Code (research preview) and a 2.5x fast mode that is three times cheaper than the previous generation.

Why did GitHub Copilot switch to token billing and what does it mean for developers?

GitHub Copilot switched from flat-rate subscriptions to token-based "AI Credits" billing on June 1, 2026 to align pricing with actual usage and address heavy users subsidized under the old model. For developers running agentic coding sessions, a single session can now cost \$30 to \$40 - far exceeding a \$10 Pro plan's monthly allotment. Microsoft's MAI-Code-1-Flash model, which uses approximately 60% fewer tokens, was launched at Build 2026 the next day to help offset credit costs for intensive users.

Is Gemini 3.5 Pro available yet?

As of June 11, 2026, Gemini 3.5 Pro has not yet shipped to general availability. Google announced the model at I/O 2026 on May 19, with CEO Sundar Pichai promising it within "next month" - meaning June 2026. It is currently in limited Vertex AI preview. Prediction markets are pricing in a late-June launch. Gemini 3.5 Flash, which did launch on May 19, already surpasses last year's Pro tier on several benchmarks and is available across AI Studio, Vertex, and the Gemini app.

What happened with DeepSeek V4 Pro pricing?

DeepSeek made its 75% price cut on V4 Pro permanent on May 23, 2026, effective from the end of the promotional period on May 31. The new permanent pricing runs from \$0.003625 to \$0.87 per million tokens - down from \$0.0145 to \$3.48. This places DeepSeek V4 Pro well below OpenAI GPT-5, Anthropic Claude Opus 4.7, and Google Gemini 3.5 Flash in cost, with a 1 million token context window. The model runs on Huawei Ascend 950 chips, reducing dependency on Nvidia hardware.

What is Anthropic's valuation for its IPO?

Anthropic's most recent private market valuation is \$965 billion post-money, set during its \$65 billion Series H funding round closed in late May 2026. This surpasses OpenAI's \$852 billion private valuation for the first time. Anthropic filed a confidential S-1 with the SEC on June 1, 2026. No IPO price, share count, or timeline has been set. The company reports a run-rate revenue of \$47 billion and is approaching its first profitable quarter.

Who are the major AI companies going public in 2026?

Three major AI-adjacent companies are pursuing public listings in 2026. SpaceX, Elon Musk's rocket and satellite company (which absorbed xAI in February 2026), is pricing its IPO on June 11 with trading starting June 12 on Nasdaq as SPCX. Anthropic confidentially filed its S-1 on June 1 at a \$965 billion valuation. OpenAI confidentially filed its S-1 on June 8 at a \$852 billion valuation. Investment bankers have advised both Anthropic and OpenAI that there is an early mover advantage to being the first AI model company to list.

Recommended Blogs

- [AI News Today - June 8, 2026: 16 Biggest Stories](#)
- [AI News Today - June 6, 2026: 16 Biggest Stories](#)
- [AI News Today - June 2, 2026: 11 Biggest Stories \(Microsoft Build\)](#)

- [AI News Today - June 1, 2026: 11 Biggest Stories](#)
- [AI News Today - May 20, 2026: Google I/O 14 Biggest Stories](#)
- [AI News Today - May 25, 2026: 15 Biggest Stories](#)
- [AI Industry News & Trends Hub - Complete BFWAI Daily Coverage](#)

Reference Links

- [OpenAI - Access OpenAI Models and Codex Through Your Oracle Cloud Commitment \(June 11, 2026\)](#)
- [OpenAI - Confidential Submission of Draft S-1 to the SEC \(June 8, 2026\)](#)
- [Reuters - ChatGPT App Hits 1 Billion Monthly Active Users in Record Time \(June 2, 2026\)](#)
- [Anthropic - Introducing Claude Opus 4.8 \(May 28, 2026\)](#)
- [Fortune - Anthropic Confidentially Files for IPO After Raising \\$65 Billion \(June 1, 2026\)](#)
- [Fortune - OpenAI Files Confidential S-1 with SEC \(June 9, 2026\)](#)
- [Capital.com - SpaceX IPO Pricing June 11, 2026 at \\$135 Per Share](#)
- [Engadget - DeepSeek Permanently Reduces V4 Pro Price by 75 Percent \(May 2026\)](#)
- [TechTimes - Google Gemini 3.5 Pro Nears June Launch with 2M Token Context \(June 6, 2026\)](#)
- [GitHub Blog - GitHub Copilot Moving to Usage-Based Billing \(June 1, 2026\)](#)
- [The Agent Report - GitHub Copilot Token Billing Backlash Hits as Microsoft Build 2026 Opens](#)
- [Build Fast With AI - AI News Today June 8, 2026](#)
- [Build Fast With AI - AI News Today June 6, 2026](#)
- [Help Net Security - Anthropic Launches Claude Opus 4.8, Prepares Mythos-Class Models](#)
- [InfoWorld - DeepSeek Steep V4-Pro Price Cut Escalates AI Pricing War](#)

Enjoyed this article? Share it →

Share:



Ship Your First AI App

From zero to deployed app with our Gen AI Launchpad

Start Building Today

0 Comments

⌵ Sort by



Sign in to join the conversation...

Ctrl+Enter to submit

📌 Sign in to comment



Be the first to comment
Start the conversation below

AI That Keeps You Ahead




Get the latest AI insights, tools, and frameworks delivered to your inbox. Join builders who stay ahead of the curve.

your@email.com

Join 50,000+ builders >

You Might Also Like



AI News

12+ AI Models in March 2026: The Week That Changed AI

Secondary Keywords GPT-5.4 release, Qwen 3.5 small benchmarks, LTX 2.3 video model, Helios ByteDance AI, NVIDIA Nemotron 3 Super





AI News

AI News Today - May 18, 2026: 13 Biggest Stories

13 AI stories you need to know today: Google I/O 2026 in 48 hrs, Anthropic \$900B round, Meta Avocado silent, GPT-5.5 Instant default, and more.

✦ Personalized Growth Engine

What's your AI Score?

Measure your AI readiness and unlock a personalized roadmap with curated tools, frameworks, and resources tailored to your role.

Check My AI Score →

✓ Takes 2 minutes

✓ Free forever

✓ Actionable advice