

[← Вернуться к блогам](#)

Новости искусственного интеллекта

AI News Today, 24 июня 2026 года: 15 главных новостей

📅 24 июня 2026 года ⌚ 21 минута чтения



Поделиться:



Компания OpenAI представила свою самую мощную на сегодняшний день модель кибербезопасности, SpaceX заключила контракт с четвертым клиентом на использование вычислительных мощностей Colossus на сумму 6,3 миллиарда долларов, а до июньского дедлайна Google по выпуску Gemini 3.5 Pro осталось шесть дней. Вот 15 событий, которые определили июнь 2026 года. Полный список событий в сфере ИИ за этот месяц

можно найти в разделе [Новости и тенденции в сфере ИИ на Build Fast with AI](#).

1. Запуск OpenAI GPT-5.5-Cyber: 85,6 % на CyberGym – самый высокий результат среди всех моделей

OpenAI запустила полную версию GPT-5.5-Cyber 22 июня 2026 года в рамках расширенной инициативы по кибербезопасности Daybreak. Модель набрала 85,6 % баллов на CyberGym (по сравнению с 81,8 % у стандартной версии GPT-5.5), 39,5 % баллов на ExploitGym (по сравнению с 25,95 %) и 69,8 % баллов на SEC-bench Pro (по сравнению с 63,1 %). OpenAI называет это самым высоким результатом на CyberGym среди всех моделей. GPT-5.5-Cyber не является общедоступной моделью API. Доступ к нему предоставляется проверенным организациям через программу Trusted Access for Cyber, в которую входят Akamai, Cisco, Cloudflare, CrowdStrike, Fortinet, Oracle, Palo Alto Networks и Zscaler.

Партнерские отношения с правительствами охватывают Австралию, Канаду, Францию, Германию, Японию, Южную Корею и институты ЕС, включая ENISA, наряду с правительством Великобритании. Модель может работать с большими кодовыми базами, отслеживать пути атак, проверять возможность использования уязвимостей, генерировать целевые исправления и предоставлять доказательства устранения уязвимостей в рамках единого автоматизированного рабочего процесса. Компания OpenAI провела тестирование перед развертыванием совместно с Центром стандартов и инноваций в области искусственного интеллекта и Управлением национального директора по кибербезопасности. Запуск напрямую связан с выполнением указа Трампа от 2 июня о безопасности искусственного интеллекта. Полный обзор конкурентной среды в сфере инструментов кибербезопасности на основе ИИ можно найти на сайте [Центр инструментов для написания кода на основе ИИ на сайте Build Fast with AI](#) рассказывает о динамике рынка Glasswing и Daybreak.

2. Patch the Planet: OpenAI, Trail of Bits и HackerOne о проблеме уязвимостей в открытом исходном коде

Alongside GPT-5.5-Cyber, OpenAI, совместно с Trail of Bits, HackerOne и Calif. запустили Patch the Planet — скоординированную инициативу по обеспечению безопасности с открытым исходным кодом. В рамках программы исследования уязвимостей с помощью искусственного интеллекта (с использованием Codex Security и GPT-5.5-Cyber) проходят обязательную проверку экспертами Trail of Bits, прежде чем информация будет передана разработчику. Такой подход является осознанным: у разработчиков с открытым исходным кодом хронически не хватает ресурсов. Согласно исследованию Linux Foundation и Гарвардского университета, в 94% широко используемых проектов менее 10 разработчиков отвечают за более чем 90% кода, добавленного за год. Если завалить разработчиков непроверенными сообщениями об ошибках, сгенерированными искусственным интеллектом, ситуация только ухудшится. Программа требует, чтобы дубликаты были удалены вручную и проверены перед публикацией. В проекте приняли участие более 30 проектов с открытым исходным кодом, в том числе cURL, Go, Python, Sigstore, pyca/cryptography, aiohttp, NATS Server, freenginx и python.org. В ходе первого пятидневного спринта были получены сотни проверенных результатов и десятки объединенных патчей. Плагин Codex Security, который с момента выхода предварительной версии в марте просканировал более 30 миллионов коммитов в 30 000 кодовых базах, напрямую интегрирован в рабочий процесс Patch the Planet.

3. SpaceX заключила с Reflection AI соглашение об аренде вычислительных мощностей на сумму 6,3 миллиарда долларов для проекта Colossus 2

22 июня 2026 года SpaceX подписала с Reflection AI договор аренды вычислительных мощностей на сумму 150 миллионов долларов в месяц, начиная с 1 июля 2026 года и до конца 2029 года, на общую сумму около 6 миллиардов долларов. 3 миллиарда, если контракт

будет действовать в течение всего срока. Reflection получает немедленный доступ к чипам Nvidia GB300 на объекте SpaceX «Колосс 2» в Мемфисе, штат Теннесси. Любая из сторон может расторгнуть контракт, уведомив об этом за 90 дней после истечения первоначального трехмесячного срока. Эта сделка стала четвертым крупным контрактом на аренду вычислительных мощностей, который SpaceX заключила для Colossus, после контрактов с Anthropic (примерно 45 миллиардов долларов до середины 2029 года на Colossus 1), Google (примерно 30 миллиардов долларов до 2029 года на Colossus 2, с ежемесячными платежами в размере 920 миллионов долларов) и Cursor. До 2029 года выручка



Панель запуска Agentic AI



крупнейших коммерческих вычислительных платформ для искусственного интеллекта в мире. Интересная структурная деталь этой сделки: Nvidia является и поставщиком чипов, и крупным инвестором в Reflection AI, создавая циклическую динамику, которую издание The Next Web охарактеризовало как 'Nvidia продает кирки и владеет частью шахты.' Ранее издание The Wall Street Journal охарактеризовало Reflection как ['DeepSeek Запада'](#) за его амбициозность в области открытых технологий и передовых разработок.

4. Reflection AI: объяснение сделки на 25 миллиардов долларов, или 'DeepSeek Запада'

Reflection AI was founded in 2024 by Misha Laskin and Ioannis Antonoglou, both veterans of Google DeepMind, where Antonoglou helped build AlphaGo. The company is backed by Nvidia, Sequoia, and Lightspeed, and raised \$2 billion in October 2025 at an \$8 billion valuation. It is now raising again at a \$25 billion valuation. The thesis: American governments, banks, defense contractors, and large enterprises want frontier-capable AI with open weights. They will not use closed US labs for sovereignty reasons, and they will not use Chinese open-weight models for security reasons. Reflection is positioning as the

third option: American, open-weight, and frontier-scale. The company has ties to the Department of Energy's Genesis Mission and Pentagon AI programs. Its code research agent Asimov has been on a waitlist since early 2026. No public frontier model has shipped yet. The \$6.3 billion compute commitment is the most concrete signal yet that a training run is imminent. Whether its 'open-weight' characterization is fully accurate remains contested: Google DeepMind engineer Susan Zhang publicly questioned whether Reflection's model release approach qualifies as genuine open source, noting that training data and the full training process remain proprietary.

5. John Jumper, Nobel Laureate and AlphaFold Lead, Joins Anthropic from DeepMind

John Jumper, who shared the 2024 Nobel Prize in Chemistry for co-leading the AlphaFold 2 team at Google DeepMind, announced on June 20, 2026 that he is leaving DeepMind after nearly nine years to join Anthropic. In a post on X, Jumper wrote that DeepMind CEO Demis Hassabis 'took a real chance letting me lead the AlphaFold team just six months after finishing my PhD, and the entire GDM team taught me so much about how to do great science.' Bloomberg reports that Jumper was also a key member of Google's coding tools team, an area where Google has struggled to convert research capability into commercial adoption. His role at Anthropic has not been officially announced, but the hire is expected to accelerate Anthropic's AI for science initiatives, including its existing work on protein structure, drug discovery, and computational biology. Anthropic already operates Project Glasswing, which has used Claude Mythos Preview to find 23,019 vulnerabilities across 1,000+ open-source projects. Jumper's scientific computing background brings a different dimension to that work. For context on Anthropic's current model lineup, the [Claude Opus 4.8 review at Build Fast with AI](#) covers benchmarks, pricing, and capabilities.

Go From AI User to AI Builder

Don't just use ChatGPT. Learn to build custom LLM agents, RAG pipelines, and full-stack Agentic AI apps in our intensive 6-week program.

- ✓ 6 Weeks Live Mentorship
- ✓ Weekly App Templates & Code
- ✓ Deploy 5+ Real-world Apps
- ✓ No Coding Experience Required

Explore Program

Join 1,000+ graduates • Free Registration

6. Google Loses Two Top AI Executives in One Week: Alphabet Stock Falls 7.2%

John Jumper's departure from DeepMind is the second high-profile AI talent loss for Google in a single week, following the reported exit of DeepMind VP John Jumper's colleague. Bloomberg reported that Alphabet stock fell as much as 7.2% intraday on the news, the steepest single-day drop since February 2026. For context, Noam Shazeer (Transformer co-author) left Google for OpenAI on June 18, and now Jumper (AlphaFold lead) has departed for Anthropic on June 20. The two moves in the same week represent a concentration of foundational AI talent leaving Google for competitors that has not happened at this pace before. The market reaction reflects investor concern about whether Google's Gemini models, which still have not delivered on the I/O commitment to launch Gemini 3.5 Pro in June, can maintain their development trajectory without these researchers. Google's \$4.46 trillion market cap gives it structural resilience, but the talent signal is real and now has a price tag attached: \$7.2% of Alphabet's intraday value on a single news cycle.

7. Gemini 3.5 Pro Still Pending GA With Six Days Left in June

As of June 24, 2026, Gemini 3.5 Pro has not achieved general availability. Google CEO Sundar Pichai committed explicitly at Google I/O on May 19 to a June 2026 launch, an audience of developers received the delay announcement with an audible groan when he pushed from I/O day to 'next month.' The model has been in limited Vertex AI enterprise preview since early June. Confirmed specifications remain: a 2-million-token context window, Deep Think reasoning mode (gated to Ultra subscribers at \$250 per month), and frontier multimodal capability. Expected pricing is \$15 per million input tokens and \$60 per million output tokens, roughly 10x the cost of Gemini 3.5 Flash. Polymarket currently prices a 50-55% probability of a June 30 GA. With six days remaining, the community is split between those expecting a late-week launch and those expecting a July slip with a formal explanation. What makes this window important beyond one model: Gemini 3.5 Pro is the only major frontier model that could reset the Gemini narrative after losing Jumper to Anthropic and Shazeer to OpenAI in the same week.

8. Gemini 3.5 Pro Leaks Show Reasoning Gap vs Fable 5 and GPT-5.6

Early leak reports from Universe of AI, cited by Geeky Gadgets, indicate that Gemini 3.5 Pro struggles in advanced reasoning, bidirectional code processing, and long-term task execution relative to Anthropic's Fable 5 and OpenAI's anticipated GPT-5.6. The specific limitations documented include weakness in code infilling and bidirectional processing, which reduce its utility for the developer coding workflows that are now the primary enterprise AI battleground. Gemini 3.5 Flash, which already shipped at I/O, improved on Gemini 3.1 Pro for coding and agentic tasks but regressed on hard reasoning. If the Pro variant has not fully closed that reasoning regression, Google's June launch will arrive with caveats. The good news for Google: Gemini 3.5 Pro's vision capabilities and multimodal understanding have reportedly improved significantly, which could make it the leader for image-heavy enterprise workflows even if it trails on pure coding and reasoning benchmarks. Also released alongside the Gemini 3.5 family is Diffusion Gemma, an experimental local-device model that processes 256-token chunks simultaneously for

up to 4x generation speed at reduced quality, which we cover separately below.

9. OpenAI Acquires Ona for Persistent Codex Sandbox Environments

OpenAI has acquired Ona, a startup that provides persistent cloud execution environments, to integrate long-running sandboxes into the Codex AI coding agent platform. The acquisition addresses one of Codex's key limitations relative to Claude Code: Codex previously operated in stateless, short-lived execution contexts. Ona's technology allows Codex to maintain persistent cloud environments for hours-to-days agent tasks, enabling the kind of long-horizon agentic work that Anthropic's Claude Code has been able to perform natively. The acquisition is a direct response to the competitive gap Anthropic has exploited to reach over 40% of the generative AI coding market. OpenAI's Codex currently holds approximately 21% of that market. For enterprise teams evaluating AI coding platforms, persistent environments are a non-negotiable requirement for any workflow that involves multi-step deployment, testing, and validation cycles spanning more than a single session.

10. Trump AI Executive Order: 30-Day Voluntary Model Review Now Active

President Trump's June 2 executive order titled 'Promoting Advanced Artificial Intelligence Innovation and Security' is now in active implementation, with OpenAI already working with the Office of the National Cyber Director and the Office of Science and Technology Policy on GPT-5.5 and GPT-5.5-Cyber testing under the new framework. The order asks AI companies to voluntarily provide the federal government with early access to covered frontier models for up to 30 days before releasing them to other trusted partners for national security cybersecurity review. It explicitly does not create mandatory government approval requirements. The original draft set the review window at 90 days; it was cut to 30 at the insistence of anti-regulation

voices within the administration, including Elon Musk and Mark Zuckerberg, who had lobbied against the earlier version. The order comes with historical context: Trump's previous AI executive order had revoked the Biden administration's AI safety order, removing mandatory pre-deployment testing requirements. The June 2 order represents a partial pivot, acknowledging frontier AI capabilities as a national security concern without restoring mandatory guardrails. The Fable 5 shutdown is the most visible evidence of what the government considers an unacceptable deployment: Anthropic's most capable models were pulled because Amazon researchers found a way through Fable 5's guardrails. For the full political context, the [Council on Foreign Relations analysis](#) covers the national security vs. innovation tradeoff in detail.

11. Diffusion Gemma: Google's 4x-Faster Local Text Generation Model Explained

Google released Diffusion Gemma alongside the Gemini 3.5 family as an experimental local-device text generation model. Instead of generating tokens sequentially, Diffusion Gemma processes 256-token chunks simultaneously, achieving up to 4x the generation speed of standard autoregressive models. The trade-off is output quality: simultaneous chunk processing sacrifices the coherence benefits of left-to-right sequential generation, making Diffusion Gemma better suited for tasks where speed matters more than precision, such as drafting, autocomplete, and rapid iteration on code comments. Features include inline editing and code infilling, making it specifically designed for developer tooling. Diffusion Gemma is intended for on-device deployment, meaning it runs locally without API calls, which has privacy and latency benefits for enterprise deployments that cannot route sensitive code through cloud APIs. It is not a replacement for Gemini 3.5 Flash or Pro; it is a speed-optimized companion for specific local use cases where a 4x inference speedup justifies reduced output quality.

12. Colossus 2 Is Now the World's Largest Commercial AI Compute Platform

SpaceX's Colossus 2 facility in Memphis, Tennessee, has become the most significant new commercial AI compute platform in the world based on committed revenue. The four signed deals including Anthropic (\$45B), Google (\$30B), Reflection AI (\$6.3B), and Cursor, combined with SpaceX's own xAI workloads, give Colossus committed external revenues exceeding \$80 billion through 2029. Colossus 2 currently houses 555,000 Nvidia GPUs purchased at approximately \$18 billion, with plans to expand the total Memphis complex to 2 gigawatts of power capacity. To put that scale in context: Amazon Web Services operates approximately 3 million servers across its global infrastructure. Colossus is a single campus approaching the compute density of a significant fraction of AWS's AI-relevant capacity in one location. The strategic implication for SpaceX: what started as an internal training cluster for Grok has become the infrastructure layer underpinning three of the four largest frontier AI labs plus a major open-weight challenger. SpaceX's compute business may generate more revenue over the next three years than its launch business.

13. MiniMax M2.5 Delivers Near-Frontier Performance at a Fraction of the Cost

MiniMax released M2.5, its latest model, with a benchmark profile that positions it as a cost-competitive alternative to frontier models from Anthropic, OpenAI, and Google. The release is part of a broader pattern in June 2026 of Chinese AI labs releasing models that approach frontier performance at significantly lower API costs. MiniMax M2.5 falls into the category of models that are good enough for most production tasks at a fraction of the cost of Fable 5, Opus 4.8, or GPT-5.5. The timing is not coincidental: MiniMax and Zhipu AI (GLM-5.2) both saw increased enterprise interest immediately following the Fable 5 suspension from June 12-18, and M2.5 positions MiniMax to capture some of that demand. For developers building multi-provider routing architectures in response to the Fable 5 experience, MiniMax M2.5 joins GLM-5.2, DeepSeek V4 Pro, and Kimi K2.7 as the primary open-weight or low-cost alternatives worth evaluating. The [best AI models June 2026 leaderboard](#) at Build Fast with AI has current benchmark comparisons across all these models.

14. ByteDance Launches Seedance 2.0 AI Video Amid Growing Platform Restrictions

ByteDance released Seedance 2.0, the latest version of its AI video generation model, with significantly improved motion consistency, longer generation windows, and enhanced character animation. The release comes as ByteDance faces intensifying regulatory scrutiny in the United States over TikTok's ownership structure. The timing creates an interesting strategic tension: ByteDance is investing heavily in AI model development that requires US compute partnerships and US developer adoption, while simultaneously navigating a political environment where its flagship consumer product remains under divestiture pressure. Seedance 2.0 is aimed at competing with Runway, Sora (OpenAI's video model), and Google's Veo in the AI video generation market. The video AI market is one of the fastest-growing segments, with enterprise adoption expanding from marketing and media production into product visualization, training data generation, and synthetic media. ByteDance's AI model capabilities are commercially strong despite the TikTok political risk.

15. SpaceX Shares Fall 16.4% as Compute Deals Raise Questions About Grok Strategy

SpaceX shares fell 16.43% on June 22, 2026, closing at \$154.60, even as the company announced the Reflection AI compute deal that adds \$6.3 billion in committed revenue through 2029. The sell-off reflects a structural concern the market is pricing: SpaceX's original AI thesis was that xAI and Grok would be a frontier AI player. The company built Colossus to train Grok. Instead, Grok has failed to gain meaningful developer market share, and SpaceX has pivoted Colossus into a compute rental business serving Anthropic, Google, Reflection, and Cursor. That pivot generates substantial revenue, but it validates the competitors rather than beating them. SpaceX is now arguably the infrastructure provider for three of the four frontier AI labs it was supposed to compete with. The market may be reassessing whether the SpaceX IPO growth story, which leaned heavily on xAI and Grok as a frontier AI moat, is

actually a compute real estate story, which carries lower multiples and different competitive dynamics. The 16% single-day drop on positive revenue news is a clear market signal that investors expected something different from SpaceX's AI strategy.

Frequently Asked Questions

Can I access GPT-5.5-Cyber through the standard OpenAI API?

No. GPT-5.5-Cyber is not available through the public OpenAI API. It is gated exclusively to organizations vetted through OpenAI's Trusted Access for Cyber program, which requires demonstrated cybersecurity use case legitimacy. Security vendors can access the model through the Daybreak Cyber Partner Program to embed it in their own products and services.

What is the Colossus 2 compute platform and how does it compare to AWS?

Colossus 2 is SpaceX's second Memphis, Tennessee data center campus, housing Nvidia GB300 chips with plans to expand the full complex to 2 gigawatts of power capacity and 555,000+ GPUs. It has signed external compute deals with Google, Reflection AI, and Cursor, alongside the original Colossus 1 lease with Anthropic. SpaceX's committed external revenues from Colossus exceed \$80 billion through 2029, making it a major commercial AI compute platform despite having no prior cloud computing business.

What research does John Jumper bring to Anthropic?

John Jumper co-led the AlphaFold team at Google DeepMind, which produced AlphaFold 2, the AI system that solved the protein-folding problem and earned the 2024 Nobel Prize in Chemistry. He brings deep expertise in applying large-scale machine learning to scientific problems, particularly in computational biology and protein structure prediction.

Anthropic has existing AI for science programs and Project Glasswing; Jumper's background is expected to accelerate scientific AI applications.

What are the benchmarks for Gemini 3.5 Flash vs Pro?

Gemini 3.5 Flash, which shipped at Google I/O on May 19, 2026, scored 76.2% on Terminal-Bench 2.1 and 83.6% on MCP Atlas, outperforming Gemini 3.1 Pro on coding and agentic benchmarks while being approximately 4x faster and significantly cheaper at \$1.50 per million input tokens and \$9 per million output tokens. Gemini 3.5 Pro is expected to close the hard reasoning gap where Flash regressed relative to 3.1 Pro, while maintaining Flash's coding and agentic improvements. Pro is priced at an estimated \$15 per million input tokens and \$60 per million output tokens.

What open-source projects are in Patch the Planet?

Over 30 open-source projects have committed to Patch the Planet, including cURL, Go, Python, Sigstore, pyca/cryptography, aiohttp, NATS Server, freenginx, and python.org. Trail of Bits has dedicated security engineers working full-time with GPT-5.5-Cyber and Codex Security across 19 of these projects. Every AI-generated finding is manually reviewed by Trail of Bits engineers before submission to a project maintainer.

Why did Alphabet stock fall 7.2% in June 2026?

Alphabet stock fell as much as 7.2% intraday following reports that John Jumper, who led the AlphaFold team and shared a Nobel Prize in Chemistry, was leaving Google DeepMind for Anthropic. The drop was amplified because it came just days after Noam Shazeer (Transformer co-author) announced his departure from Google for OpenAI on June 18. Two foundational AI researchers leaving Google for competitors in the same week, while Gemini 3.5 Pro remains undelivered, created compounding investor concern about Google's AI talent retention and execution.

What is Diffusion Gemma and who should use it?

Diffusion Gemma is an experimental local-device text generation model from Google that processes 256-token chunks simultaneously instead of sequentially, achieving up to 4x generation speed at reduced output quality. It supports inline editing and code infilling. It is designed for developers who need rapid local inference, such as autocomplete, drafting, and code comment generation, where speed matters more than precision. It is not a replacement for Gemini 3.5 Flash or Pro.

What is MiniMax M2.5 and how does it compare to GLM-5.2?

MiniMax M2.5 is a near-frontier Chinese AI model released in June 2026, positioned as a cost-competitive alternative to frontier models from OpenAI, Anthropic, and Google. Like GLM-5.2 from Zhipu AI, M2.5 targets developers and enterprises who need strong performance at lower API costs than Fable 5 or GPT-5.5. Both M2.5 and GLM-5.2 have benefited from increased enterprise interest following the Fable 5 suspension, as organizations evaluate multi-provider fallback architectures.

Recommended Blogs

- [AI News Today June 23 2026: Fable 5 Billing Cliff, SpaceX Cursor, Gemini 3.5 Pro Window](#)
- [AI News Today June 22 2026: Fable 5 Free Trial Ends, Noam Shazeer Joins OpenAI](#)
- [Best AI Models June 2026: Full Ranked Leaderboard](#)
- [Claude Code Review: Benchmarks, Pricing, and Agentic Capabilities](#)
- [AI News Today June 2 2026: Microsoft Build, Anthropic IPO Filing](#)
- [GLM-5.2 Review: Open-Weight Frontier Model at 1/6th the Cost](#)

Resources and Community

Join our community of **70,000+** AI enthusiasts and learn to build powerful AI applications! Whether you are a beginner or an experienced developer, Build Fast with AI helps you understand and implement AI in your projects.

- [Website — buildfastwithai.com](https://buildfastwithai.com)
- [LinkedIn — Build Fast with AI](#)
- [Instagram — @buildfastwithai](#)
- [Founder Twitter — @satvikps](#)
- [Twitter — @BuildFastWithAI](#)

Agentic AI Launchpad 2026

A structured **6-week cohort program** that takes you from AI basics to building and deploying real-world agentic AI systems. Includes live sessions, expert mentorship, project reviews, and a builder community network.

Ready to go from learning to building? Join the next cohort: [Agentic AI Launchpad 2026](#)

Free AI Resources

Access free tools, workshops, and micro-learning to keep building:

- [AI Workshops — Free resources, upcoming events and past recordings](#)
- [Unrot — Learn AI in 5 minutes a day \(free micro-learning app\)](#)

The AI news cycle is compressing. Subscribe to the Build Fast with AI newsletter and follow @BuildFastWithAI on X to get sourced daily AI coverage before the morning brief.

References

- [CybersecurityNews — OpenAI Releases GPT-5.5-Cyber](#)
- [Axios — OpenAI Rolls Out More Capable Version](#)
- [Digital Applied — GPT-5.5-Cyber and Codex Security](#)
- [The Decoder — OpenAI Says GPT-5.5-Cyber](#)
- [The Next Web — SpaceX Signs \\$6.3B Compute Deal](#)
- [Yahoo Finance — SpaceX Signs \\$6.3 Billion Compute](#)
- [MLQ News — SpaceX Signs \\$6.3B Compute Deal](#)
- [andrew.ooo — What Is Reflection AI: The \\$25B Open-Source](#)
- [TechCrunch — Nobel Laureate John Jumper Is](#)
- [Geeky Gadgets — Google Gemini 3.5 Pro Leaks Show Struggles Against GPT-5.6](#)
- [TechTimes — Google Gemini 3.5 Pro Nears June Launch](#)
- [CFR — Assessing Trump's Executive Order on AI Oversight](#)
- [CNBC — Trump Signs AI Executive Order Asking](#)
- [Cryptopolitan — Reflection AI Signs \\$6.3 Billion Compute](#)
- [Build Fast with AI — Best AI Models June 2026 Full Leaderboard](#)

Enjoyed this article? Share it →

Share:



Ship Your First AI App

From zero to deployed app with our Agentic AI Launchpad

Start Building Today

0 Comments

Sort by



Sign in to join the conversation...

Ctrl+Enter to submit

Sign in to comment



Be the first to comment
Start the conversation below

AI That Keeps You Ahead



Get the latest AI insights, tools, and frameworks delivered to your inbox. Join builders who stay ahead of the curve.

your@email.com

Join 50,000+ builders

You Might Also Like






AI News

AI News Today - June 8, 2026: 16 Biggest Stories

16 AI stories: Apple WWDC 2026 recap, Claude becomes an iPhone option, Microsoft Foundry, EU AI Act countdown, Pentagon AI race, SpaceX IPO, and more.



AI News

12+ AI Models in March 2026: The Week That Changed AI

Secondary Keywords GPT-5.4 release, Qwen 3.5 small benchmarks, LTX 2.3 video model, Helios ByteDance AI, NVIDIA Nemotron 3 Super

 Personalized Growth Engine

What's your **AI Score**?

Measure your AI readiness and unlock a personalized roadmap with curated tools, frameworks, and resources tailored to your role.

Check My AI Score →

✓ Takes 2 minutes ✓ Free forever ✓ Actionable advice