

Anthropic заявляет, что Alibaba должна понести наказание за крупнейшую атаку по клонированию Claude

Alibaba предположительно использовала 25 000 аккаунтов для майнинга Claude на 28,8 миллиона биржах.

ЭШЛИ БЕЛАНЖЕ – 25 ИЮНЯ 2026 Г. 21:01 | 142



Источник: NurPhoto / Contributor | NurPhoto

НАСТРОЙКИ ТЕКСТА



Компания Anthropic обвинила китайскую фирму Alibaba в крупнейшей на сегодняшний день попытке клонировать Claude. Китай стремится сравняться по возможностям с ведущей моделью Anthropic после выхода Mythos и последующего ограничения доступа к зарубежным рынкам.

Издание Ars получило письмо от 10 июня, отправленное сенаторам Тиму Скотту (республиканец от штата Южная Каролина) и Элизабет Уоррен (демократ от штата Массачусетс) за день до слушаний в сенатском комитете на тему «Искусственный интеллект и американская мечта». В письме компания Anthropic поделилась «новыми конфиденциальными доказательствами крупнейшей кампании по незаконному использованию возможностей Claude, которую мы когда-либо отслеживали».

Атаки происходили в период с 22 апреля по 5 июня, когда «операторы, связанные с Alibaba и Alibaba Qwen, лабораторией искусственного интеллекта Alibaba», предположительно провели «более 28,8 миллиона транзакций с Claude через почти 25 000 мошеннических аккаунтов», сообщили в Anthropic. Эта кампания, нарушающая условия предоставления услуг и ограничения доступа к Claude, «была нацелена на некоторые из наиболее ценных возможностей Claude, таких как агентное мышление, разработка программного обеспечения и долгосрочные задачи».

По данным Anthropic, Alibaba избежала обнаружения, «используя методы обфускации и прокси-сети». Anthropic предупреждает, что по мере роста спроса на надежные методы обфускации в Китае будет расширяться «экономика обхода», что приведет к еще большему количеству атак с дистилляцией в будущем.

Alibaba предположительно проигнорировала предупреждение Трампа

Как и другие китайские лаборатории, пытающиеся скопировать передовые американские модели, целью Alibaba, по утверждению Anthropic, было перенять возможности Claude, «не неся при этом затрат на обучение и научно-исследовательские и опытно-конструкторские работы, необходимые для обучения» собственной передовой модели. Такие атаки стали «широко распространенными» и «превращают сотни миллиардов долларов американских инвестиций и научно-исследовательских и опытно-конструкторских работ в масштабное субсидирование наших геополитических конкурентов», — заявили в Anthropic.

ARS VIDEO

How Lighting Design In The Callisto Protocol Elevates The Horror



Важно отметить, что, по данным Anthropic, кампания Alibaba была проведена после того, как Дональд Трамп предпринял шаги по пресечению подобных незаконных атак и защите национальной безопасности США. В апреле Трамп обвинил Китай в «промышленном» воровстве технологий искусственного интеллекта после того, как Anthropic обвинил китайские компании DeepSeek, Moonshot и MiniMax в использовании той же тактики, что и Alibaba, для создания «более 16 миллионов обменов с Claude через примерно 24 000 мошеннических аккаунтов». По данным Anthropic, OpenAI и Google опубликовали результаты расследования аналогичных атак на их модели.

ADVERTISEMENT

Компания Anthropic обвинила Alibaba в «наглой» попытке создать клонированный аналог Claude, не обращая внимания на угрозы Трампа принять меры против попыток иностранных компаний копировать американские передовые модели, несмотря на то, что Alibaba зависит от американских инвесторов.

«Alibaba зарегистрирована на Нью-Йоркской фондовой бирже, ведет бизнес в США и подотчетна американским инвесторам и регулирующим органам, — отмечается в письме Anthropic. — Однако эта деятельность развернулась спустя несколько недель после того, как» в меморандуме Трампа было заявлено, что попытки клонирования «неприемлемы».

Ars could not immediately reach Alibaba for comment.

Anthropic wants firms like Alibaba punished

Alibaba is already preparing to clash with Trump, though. In a lawsuit filed Tuesday, Alibaba accused the Trump administration of blacklisting the company after falsely linking the company to the Chinese military, Reuters reported. Alibaba is seeking to remove the Trump designation, which they claimed has “no basis in fact or law.”

“Alibaba is governed by an independent board, none of whom has any military affiliation,” Alibaba said. “Its products and services are built for retail, logistics, and enterprise information technology—not weapons, defense, or intelligence.”

Anthropic appears unconvinced, however, that Alibaba isn’t working with the Chinese government. In the letter, Anthropic warned that without stronger interventions, these distillation attacks will “help China reach Mythos Preview-level capabilities sooner.”

To keep the US ahead of China, Anthropic recommended that Congress pass legislation with three objectives. First, antitrust laws must be updated to allow AI firms to share information about evolving Chinese tactics to deter more threats.

Second, the US needs more export controls on chips to hamstring Chinese access to advanced compute so that they simply can’t train on US model outputs. That could make conducting distillation attacks pointless, Anthropic suggested.

Finally, Congress should pass laws penalizing Chinese labs’ “bad behavior” so that it’s “more difficult and costly” to rely on distillation attacks to advance Chinese models. Penalties could include limiting Chinese firms from accessing US models or advanced US chips or from relying on data centers outside of China, Anthropic suggested.

ADVERTISEMENT

Anthropic declined to clarify whether Alibaba’s alleged attacks were significant enough to help meaningfully accelerate China’s AI capabilities or comment on any specific steps taken to thwart the attacks. Instead, a spokesperson provided a statement to Ars, echoing sentiments expressed in the letter to senators.

“We believe combating the threat of illicit distillation requires coordinated action between government and industry, and we will continue working with Congress and the Administration to maintain American AI leadership,” Anthropic said.

China races to match Mythos’ capabilities

Anthropic’s letter positions the AI firm as intent on helping the US hold the line so that China cannot surpass US capabilities.

If that happened, Anthropic warned that China could blindside a defenseless US—suddenly possessing “advanced cyber capabilities to deploy against the US government and American companies and exploit vulnerabilities faster than previously possible.”

It’s important to keep the US as far ahead as possible, Anthropic’s letter said, because “the larger the capability gap,” the “more time the US government will have to harden cyber defenses and adopt AI systems across national security domains” as China’s AI advances.

Additionally, Anthropic warned that if the US ignores distillation attacks, China could release advanced AI models “with weak safeguards that are easily jailbroken, enabling other US adversaries to use these models for a wide range of activities that run contrary to US interests.”

Alibaba’s models have been downloaded more than 700 million times and are at the frontier of China’s AI industry. The official newspaper of the Central Committee of the Communist Party of China (CPC), People’s Daily, recently [hyped](#) Alibaba’s Qwen family of AI models as “the most popular open-source AI system worldwide.” The AI firm will likely maintain a defensive posture as US scrutiny escalates, but the company risks hobbling its business the longer its US fights endure. Alibaba’s stock dropped 3 percent after Anthropic’s accusations became public, Yahoo Finance [reported](#).

ADVERTISEMENT

Anthropic’s suspicions that China is racing to build models to match Claude’s capabilities have been confirmed by at least one major Chinese tech founder. At a cybersecurity conference in Beijing yesterday, 360 Security Technology founder Zhou Hongyi likened Anthropic’s Mythos to a “cyber nuclear weapon,” the South China Morning Post [reported](#).

Zhou told the audience that Mythos' sudden giant leap in its ability to find cybersecurity vulnerabilities was a “terrifying change” that had effectively “democratized” cyberattacks, SCMP reported.

For China, having no access to Mythos was a significant disadvantage, Zhou said. He bemoaned that Project Glasswing, which granted more than 40 US organizations access to Mythos Preview to strengthen cyber defenses, excluded China.

“This means US organizations can use Mythos to scan your vulnerabilities, but you don’t even have the qualification to catch a glimpse of Mythos,” Zhou said.

China’s only way forward is to create its own Mythos-like model, Zhou said, warning that such a “game-changing weapon in cyber warfare” cannot “be held solely in the hands of others.” According to Zhou, China must race to copy Mythos’ capabilities so that there’s mutually assured destruction should its rival attempt to seize gains using its advanced AI.

SCMP noted that “Zhou’s remarks marked the first time a prominent Chinese technology founder has publicly warned about the strategic risks posed by the US frontier AI model.”

Right now, Zhou said that Chinese firms are “well short of Mythos-level capabilities,” SCMP reported. He then positioned his own company as developing a solution, which focuses “on AI agent systems that combined existing foundation models with specialist security data sets and vulnerability knowledge bases,” instead of “trying to match the US in frontier model capability and computing power.”



ASHLEY BELANGER SENIOR POLICY REPORTER

Ashley is a senior policy reporter for Ars Technica, dedicated to tracking social impacts of emerging policies and new technologies. She is a Chicago-based journalist with 20 years of experience.

 **142 COMMENTS**

 **PREV STORY**

NEXT STORY 

 **MOST READ**



- 1. Anthropic says Alibaba must be punished for largest Claude cloning attack**
- 2. Feds deny Polestar authorization to sell cars in US from model year 2027**
- 3. Microsoft adds another year to Windows 10 extended update program**
- 4. IBM claims world's first sub-1 nanometer chip technology**
- 5. Underpromise, overdeliver? Hands-on with the \$24,950 Slate auto.**



Ars Technica has been separating the signal from the noise for over 25 years. With our unique combination of technical savvy and wide-ranging interest in the technological arts and sciences, Ars is

the trusted source in a sea of information. After all, you don't need to know everything, only what's important.



MORE FROM ARS

[ABOUT US](#)

[STAFF DIRECTORY](#)

[ARS NEWSLETTERS](#)

[GENERAL FAQ](#)

[POSTING GUIDELINES](#)

[AI POLICY](#)

[RSS FEEDS](#)

CONTACT

[CONTACT US](#)

[ADVERTISE WITH US](#)

[REPRINTS](#)



Manage Preferences

© 2026 Condé Nast. All rights reserved. Use of and/or registration on any portion of this site constitutes acceptance of our [User Agreement](#) and [Privacy Policy and Cookie Statement](#) and [Ars Technica Addendum](#) and [Your California Privacy Rights](#). Ars Technica may earn compensation on sales from links on this site. [Read our affiliate link policy](#). The material on this site may not be reproduced, distributed, transmitted, cached or otherwise used, except with the prior written permission of Condé Nast. [Ad Choices](#)