


ИИ / РАЗРАБОТКА ПРОГРАММНОГО ОБЕСПЕЧЕНИЯ

Советы разработчикам по созданию подсказок для ИИ

Бесплатный курс OpenAI по пром프트-инжинирингу рассказывает о том, как разработчики могут улучшить код, генерируемый ИИ, и дает советы по составлению промтов.

25 мая 2023 года, 10:35 [Лорейн Лоусон](#)



Изображение с помощью глубокого обучения.ИИ

Диапазон зарплат в сфере пром프트-инжиниринга составляет от 212 000 до 335 000 долларов за работу, которая, по сути, требует уверенного владения английским языком, — по крайней мере, так вам внушают **при чтении об этой новой технологической профессии**.

Но я обнаружил, что все немного сложнее, когда проходил курс DeepLearning.AI и **бесплатный курс OpenAI по пром프트-инжинирингу для разработчиков**. Курс показал, что навыки разработчика, знание Python и **Ytter Notebook** — ключевые факторы для тонкой настройки пром프트-инжиниринга и получения пригодного для использования кода.

приложений, все еще сильно недооценен», — сказал преподаватель курса Эндрю Ын, основатель DeepLearning.AI, соучредитель Coursera и адъюнкт-профессор факультета компьютерных наук Стэнфордского университета. «На самом деле моя команда в AI Fund, дочерней компании DeepLearning.AI, работает со многими стартапами над применением этих технологий в самых разных сферах». Было интересно наблюдать за тем, какие API позволяют разработчикам очень быстро создавать продукты».

Изабелла Фулфорд, сотрудница технического отдела OpenAI, присоединилась к Нгу в качестве со-инструктора на часовом курсе (время на самостоятельную работу может варьироваться). В курсе проводится различие между базовыми языковыми моделями и языковыми моделями, обученными на инструкциях. Языковые модели, обученные на инструкциях, предсказывают следующее слово на основе текстовых обучающих данных и часто обучаются на больших массивах данных из интернета и других источников, чтобы определить наиболее вероятное следующее слово. По словам Нгу, большие языковые модели, обученные следовать инструкциям и отвечать на вопросы, стали основой для многих исследований и практик в этой области.

По словам Нгу, LLM, обученные на инструкциях (основная тема этого курса), нацелены на то, чтобы быть полезными, честными и безобидными, поэтому вероятность того, что они выдадут проблемный текст, например токсичный, ниже, чем у базовых LLM.

First Principle: Specific Is Better

The first principle for AI development Ng explored was how to give instructions to an LLM.

“When you use an instruction-tuned LLM, think of giving instructions to another person — say someone who’s smart, but doesn’t know the specifics of your task,” Ng said. “When an LLM doesn’t work, sometimes it’s because the instructions weren’t clear enough. For example, if you were to say, please write me something about Alan Turing. Well, in addition to that, it can be helpful to be clear about whether you want the text to focus on his scientific work or his personal life or his role in history or something else.”

1. OpenClaw and Hermes agree on what an agent is. They disagree on what controls it.
2. MCP gets its missing enterprise authorization layer
3. Developers are now validating code they didn't write — and may not understand
4. Can DNS become the basis for AI agent identity?
5. AI agents need infrastructure: Why Europe's regional cloud strategy matters

It also helps to specify what tone you want the answer in, he said. You might want a professional journalism tone or prefer a more casual tone for a friend. You might also input and specify any text snippets the AI should leverage to create its draft.

“You should express what you want a model to do by providing instructions that are as clear and specific as you can possibly make them,” Fulford said. “This will guide the model towards the desired output and reduce the chance that you get a relevant or incorrect response.”

Clear writing doesn't necessarily mean creating a short prompt, as in many cases longer prompts actually provide more clarity and context for the model, leading to more detailed and relevant outputs, she added. She outlined several tactics to create specific prompts.

“The first tactic to help you write clear and specific instructions is to use delimiters to clearly indicate distinct parts of the input,” Fulford said, pasting an example into Jupyter Notebooks. “Delimiters can be any clear punctuation that separates specific pieces of text from the rest of the prompt — these could be backticks, you could use quotes, you could use **XML** tags, section titles, anything that makes it clear to the model that this is a separate section,” she said. “Using delimiters is also a helpful technique to try and avoid prompt injections; and what a prompt injection is, is if a user is allowed to add some input into your prompt, they might give conflicting instructions to the model that might make it follow the user's instructions, rather than doing what you wanted it to do.”

“The next tactic is to ask for a structured output,” she continued. “To make parsing the model up, it can be helpful to ask for a structured output like HTML or JSON.”

model to check these assumptions first, and then if they're not satisfied, indicate this and stop short of a full task completion attempt," she said.

"Our final tactic for this principle is what we call few shot prompting, and this is just providing examples of successful executions of the task you want performed before asking the model to do the actual task you want it to do," she said.

Second Principle: Give the Model Time to 'Think'

It's also important to give the LLM time to "think," Fulford said.

"If a model is making reasoning errors by rushing to an incorrect conclusion, you should try reframing the query to request a chain or series of relevant reasoning before the model provides its final answer," she said. "Another way to think about this is that if you give a model a task that's too complex for it to do in a short amount of time, or in a small number of words, it may make up a guess which is likely to be incorrect."

This error would also happen to a person if you asked a complex math question without time to work out the answer, she said — they would likely make a mistake.

"In these situations, you can instruct the model to think longer about a problem, which means it's spending more computational effort on the task," she said. For example, she asked the model to determine if a student's answer to a math problem was correct and it determined it was — except that the answer wasn't right.

A better approach, she explained, was to have the model itself solve the problem and then compare it with the student's solution, an approach that indeed revealed the student had the incorrect answer.

"The model just has agreed with the student because it skim-read it in the same way that I just did," Fulford said. "We can fix this by instructing the model to work out its own solution first, and then compare its solution to the student's solution."

That obviously requires a longer, more involved prompt than simply asking if the student's answer is correct.

re tactics for giving the model time to think:

Reducing Hallucinations

One LLM limitation is hallucinations, which is basically when the AI makes up something that sounds plausible but isn't actually correct.

"Even though the language model has been exposed to a vast amount of knowledge during its training process, it has not perfectly memorized the information [...] and so it doesn't know the boundary of its knowledge very well," Fulford said. "This means that it might try to answer questions about obscure topics and can make things up that sound plausible but are not actually true."

One way to reduce hallucinations is to ask the model to first find any relevant quotes from the text and then ask it to use those quotes to answer questions, she said. "Having a way to trace the answer back to [a] source document is often pretty helpful to reduce these hallucinations," she added.

These are just a few key principles from the course, which also explores iterative prompt development; summarizing text a focus on specific topics; inferring sentiment and topics from product reviews and news articles; and transformation tasks such as language translation, spelling and grammar checking, tone adjustment and format conversions. The final lesson shows how to create a chatbot with OpenAI.

Each video lesson is accompanied by examples and prompts that developers can try out and tinker with on their own. There's also a community section that you can join for further help.

TNS



Loraine Lawson is a veteran technology reporter who has covered technology issues from data integration to security for 25 years. Before joining The New Stack, she served as the editor of the banking technology site Bank Automation News. She has...

[Read more from Loraine Lawson →](#)

TRENDING STORIES

1. OpenClaw and Hermes agree on what an agent is. They disagree on what controls it.
2. MCP gets its missing enterprise authorization layer
3. Developers are now validating code they didn't write — and may not understand
4. Can DNS become the basis for AI agent identity?
5. AI agents need infrastructure: Why Europe's regional cloud strategy matters

INSIGHTS FROM OUR SPONSORS



VMware Tanzu

USSFCU Secures Financial Data on a Sovereign Private Cloud

25 June 2026

July 22, 2026: Demystifying the Agentic Loop: Building Safe, Server-Driven Agents with Spring

22 June 2026

July 8, 2026: Making the Move to Real-Time Financial Data Intelligence with VMware Tanzu

18 June 2026



Verification Is Not Testing. Understanding the Difference Could Save Your Codebase

June 2026

The AI Code Verification Bottleneck: Why Faster Code Generation Means Slower Reviews

3 June 2026



Your Vercel AI SDK app is missing a session layer

25 June 2026

Stop vs disconnect - why canceling AI streaming is harder than it looks

25 June 2026

Agentic apps that go beyond chat

24 June 2026

TNS DAILY NEWSLETTER

Receive a free roundup of the
most recent TNS articles in
your inbox each day.

svo1975pnv1982@gmail.com

SUBSCRIBE

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions

ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds



Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap



© The New Stack 2026

[Disclosures](#)

[Terms of Use](#)

[Advertising Terms & Conditions](#)

[Privacy Policy](#)

[Cookie Policy](#)