

ИИ / РАЗРАБОТКА ИИ / DEVOPS / РАЗРАБОТКА ПРОГРАММНОГО ОБЕСПЕЧЕНИЯ

Компания Qodo только что представила функцию проверки кода в разных репозиториях. Вот почему это важно для команд, активно использующих искусственный интеллект.

Искусственный интеллект заваливает команды запросами на слияние, которые слишком велики для проверки человеком. Новые инструменты Qodo изучают ваши стандарты кодирования и выявляют ошибки в разных репозиториях до слияния.

22 июня 2026 года, 15:42 [Адриан Бриджуотер](#)




Фото МАРИОЛЫ ГРОБЕЛЬСКОЙ на Unsplash

Ночные сеансы отладки — это уже не редкий случай. Они происходят, когда процесс проверки не поспевает за объемом работы. Именно поэтому процесс

Итамар Фридман, соучредитель и генеральный директор компании Qodo, рассказывает *The New Stack*, что монолитные стеки приложений уступили место инфраструктуре с несколькими репозиториями для современных корпоративных приложений. И последствия нестабильности и уязвимости приложений, возникающие из-за взаимодействия репозитория, сильнее всего сказываются на людях, а не на конвейере.

По его словам, теперь инженер-программист может потратить два дня на то, чтобы «разделить на две части транзитивный сбой», который нарушил работу нижестоящего сервиса.

«Изменение в одну строку в одном репозитории незаметно нарушает архитектурный инвариант (условие, которое всегда выполняется в определенной точке программы), на который полагалась другая команда, и никто этого не заметил, потому что рецензент бегло просматривал список изменений в коде из 500 строк, а его собственный дедлайн приближался», — говорит Фридман.

По мере того как ИИ генерирует все больше и больше пул-реквестов, Фридман предполагает, что «радиус поражения» будет только расти.

Поскольку агенты на основе ИИ коренным образом изменили подход к созданию программного обеспечения, системы управления, которые компании разрабатывали для разработки в условиях человеческого фактора, не были рассчитаны на агентный жизненный цикл разработки программного обеспечения. Согласно отчету «**Google DORA 2025** Состояние разработки программного обеспечения с использованием ИИ», пул-реквесты от команд с высоким уровнем внедрения ИИ уже на 154% больше по размеру, на их проверку уходит на 91% больше времени, а в релизе на 9% больше ошибок.

Код, сгенерированный искусственным интеллектом, в корпоративном масштабе

Qodo хочет дать командам разработчиков программного обеспечения возможность сохранять контроль в таких условиях. В понедельник компания анонсировала три новые возможности платформы: проверку кода в разных репозиториях, поиск пользовательских правил и стандарты проверки языков. Эти новые возможности устраняют ряд проблем с управлением,

ТРЕНДОВЫЕ ИСТОРИИ

1. OpenClaw and Hermes agree on what an agent is. They disagree on what controls it.
2. MCP gets its missing enterprise authorization layer
3. Developers are now validating code they didn't write — and may not understand
4. Can DNS become the basis for AI agent identity?
5. AI agents need infrastructure: Why Europe's regional cloud strategy matters

Multi-repo advocates talk about the separation of concerns made possible by disaggregating software based on different functional domains (authentication in one place, search in another, etc.), different stack disciplines (database drivers in one place, user interface controls in another), or quite simply a separation based upon different teams' ownership status.

Useful for isolating changes and sometimes helpful for compliance reasons, a more heterogeneous code environment also comes with a cost in complexity, reduced navigability, and more onerous maintenance. With AI now creating its own new streams of code, multi-repo setups are likely to become more prevalent in the immediate future.

The mechanics in motion here mean that as software engineering teams grow, the most important and consequential bugs are rarely contained within a single repository; more often, they have tentacles in several places. A change to a shared library, an exported API, a data schema, or an infrastructure file can introduce breaking changes across dozens of downstream services, with no warning surfaced at the point of merge.

Because coding standards often reside in wikis, developer annotations and comments, and the institutional memory of senior engineers, Qodo Custom Rules Miner is designed to disrupt the status quo. Rather than requiring teams to define standards before enforcement can begin, it automatically discovers coding

Standards in people's heads, "subjective tribal knowledge"

"These kinds of standards were previously unenforceable, all because they lived in people's heads or across system boundaries that no single tool could see: cross-repository architectural invariants, the subjective tribal knowledge of senior engineers, and AI agent workflows running outside any organizational visibility or control," explains Friedman

"[One team] gamed the metric perfectly while testing nothing. That's tribal knowledge failing twice over: the real standard (tests must assert behavior) lived in one senior engineer's head and was never encoded, so the organization optimized the proxy instead of the thing."

Qodo's Skill Review Standards service is designed to formalize and simplify standardization procedures. When software teams use agent skills to encode development workflows and best practices, they also need to manage those skills, creating a governance challenge of its own. Qodo now provides centralized management for skills that contain code review instructions, coding standards, and engineering best practices. The platform discovers skills across repositories, surfaces them in a dedicated portal experience, and enables teams to control and measure their impact.

How bad can a slip in code quality get?

Friedman shares a personal favorite example of code quality checking gone wrong in an anecdote related to code coverage, i.e., the measure of how much of a developer's codebase is executed when automated tests are run, which many would consider a proxy or even a vanity metric.

A friend of his needed to reach 80% code coverage. One day, he checked one of the repos, and they had achieved 85% coverage; "he was so proud", until he noticed there were no asserts (an assertion is an automated verification test that checks whether code produced the result expected) at all.

(tests must assert behavior) lived in one senior engineer's head and was never encoded, so the organization optimized the proxy instead of the thing. That's precisely the gap the rules system in Qodo closes, by learning the real patterns from how good engineers work and making them enforceable."

"Speed without independent verification isn't velocity; it's technical debt in disguise."

What developers should think next

There's a lot to change here; it's a question of new tools, new workflows (re-engineered human-coded ones and AI-generated ones), and new workplace practices. We're at an inflection point, seeing a shift from stateless AI tools to stateful systems with persistent organizational memory.

Friedman suggests that we can capitalize upon this. He says that as the volume of AI-generated code quickly outpaces human review, governance systems that can learn, remember, and enforce organizational standards are becoming a near-term necessity rather than a future aspiration.

"Cross-repo review and centralized agent management matter, but without memory, they're still stateless. You can't be good at code review without a very strong context engine underneath it, and memory is that engine," he surmises.

The bottom line from Qodo is that developers should realize that they can't trust the same generalist model that wrote the code to objectively grade its own work. We already know that LLMs will confidently tell us that code works when it does not.

Friedman's bid for a t-shirt slogan here is: Speed without independent verification isn't velocity; it's technical debt in disguise. As the volume of AI-generated code has already outpaced human-paced review, the challenge is to embed governance into the development system itself through machine-readable standards, verification loops, and cross-system visibility, so that humans can oversee the process rather than inspect every line of code.



television. Primarily working as an analysis writer dedicated to a software application development 'beat',...

[Read more from Adrian Bridgwater →](#)

SHARE THIS STORY



TRENDING STORIES

1. OpenClaw and Hermes agree on what an agent is. They disagree on what controls it.
2. MCP gets its missing enterprise authorization layer
3. Developers are now validating code they didn't write — and may not understand
4. Can DNS become the basis for AI agent identity?
5. AI agents need infrastructure: Why Europe's regional cloud strategy matters

INSIGHTS FROM OUR SPONSORS



Open Governance for MySQL: A Step Forward for the Community

25 June 2026

Governing AI Assets at Scale with MCP Gateway and Registry

17 June 2026



SUSE

SUSE Multi-Linux Manager 5.x on Azure

24 June 2026

Running highly available virtualization at the Edge with just two data nodes

24 June 2026

Navigating the SAP Tightrope: Innovation vs. Control in a Volatile World

23 June 2026



Red Hat
OpenShift

Friday Five — June 26, 2026

26 June 2026

From alert fatigue to automated action: Automated patching in the AI era

25 June 2026

Govern privileged workload boundaries with Red Hat OpenShift, Ansible Automation Platform, and Identity Management

25 June 2026

TNS DAILY NEWSLETTER

Receive a free roundup of the
most recent TNS articles in

svo1975pnv1982@gmail.com

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions

FOLLOW TNS



ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds



Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

[Disclosures](#)

[Terms of Use](#)

[Advertising Terms & Conditions](#)

[Privacy Policy](#)

[Cookie Policy](#)