


ИНФРАСТРУКТУРА ИИ / ДАННЫЕ / БАЗЫ ДАННЫХ

Проблема с хранением данных в базе решена. Вот что будет дальше.

Postgres превращается из системы хранения данных в центр их перемещения. Вот как интероперабельность и искусственный интеллект меняют его будущее.

18 июня 2026 года, 10:00 [Крейг Керстиенс](#)



Эдуардо Рамос для Unsplash+

[Snowflake](#) спонсировала этот пост.

На протяжении большей части своей 30-летней истории Postgres считалась транзакционной базой данных. Организации доверяют ей хранение данных о клиентах, финансовых транзакциях и выполнение множества других операционных задач. Ее репутация основана на надежности, строгих гарантиях транзакций и активном сообществе разработчиков с открытым исходным кодом, которое десятилетиями совершенствовало базу данных, не нарушая ее фундаментальных принципов.

сокращением необходимости в их перемещении.

«Некоторые из самых важных инноваций в экосистеме Postgres сегодня связаны не столько с хранением данных, сколько с сокращением необходимости их перемещения.»

Инновации в области баз данных исторически были направлены на повышение производительности, масштабируемости и надежности. Все более сложной проблемой становится интероперабельность: как обеспечить обмен операционными данными между аналитическими системами, приложениями с искусственным интеллектом и последующими сервисами без создания еще одного конвейера или копии данных.

Почему Postgres продолжает появляться в наших списках

Реальность современной архитектуры программного обеспечения такова, что данные редко хранятся в одном месте. Информация, созданная в операционных системах, быстро попадает на склады, поисковые платформы, в среды машинного обучения и приложения с искусственным интеллектом. Каждая новая система решает актуальную бизнес-задачу, но при этом создает еще одно хранилище для данных и зачастую еще одну копию, которую нужно поддерживать.

Затраты при таком подходе не ограничиваются расходами на инфраструктуру. Каждая дополнительная копия приводит к задержкам, создает еще один потенциальный источник несогласованности и увеличивает операционную нагрузку, связанную с синхронизацией систем. Многие организации тратят на перемещение данных столько же усилий, сколько на их хранение.



Snowflake позволяет любой организации мобилизовать свои данные с помощью Data Cloud от Snowflake. Клиенты используют Data Cloud для объединения разрозненных данных, поиска и безопасного обмена данными, создания

[Learn More →](#)

HEAR MORE FROM OUR SPONSOR

svo1975pvn1982@gmail.com[Submit](#)

“Many organizations now spend as much effort moving data as they do storing it.”

For many businesses, Postgres serves as the system of record for customer interactions, transactions, application state, and other business-critical information. As organizations expand their analytical, machine learning, and AI capabilities, they are not looking to create another source of truth; rather, they're looking for better ways to work with the one they already trust.

That shift is changing how Postgres fits into modern architecture. Historically, Postgres was viewed primarily as the place where operational data originated before being copied into downstream systems. Increasingly, organizations want those systems to work more seamlessly with operational data while reducing the pipelines, copies, and synchronization processes required to support them.

TRENDING STORIES

1. [The database storage problem is solved. Here's what comes next.](#)
2. [When your data model is the bottleneck: lessons from Medium's feature store](#)
3. [Why MotherDuck refuses to fork DuckDB](#)
4. [Why the need for humans won't disappear in the age of autonomous databases](#)

Technologies such as **logical replication**, change data capture, and foreign data wrappers have helped Postgres participate more directly in larger data ecosystems. As a result, organizations are no longer asking only whether Postgres can store their data. They're instead asking how easily it can connect to everything around it.

That shift, from evaluating databases primarily on storage and performance to evaluating them on interoperability, may be one of the most important changes happening in the Postgres ecosystem today.

AI is exposing old problems

The recent focus on AI has brought renewed attention to data movement. AI didn't create the problem. If anything, it exposed a limitation that has been quietly growing for years. For decades, organizations built architectures around the idea that data would move between systems through pipelines and periodic synchronization. That model worked because most analytical workloads could tolerate some degree of delay.

AI is changing those expectations. Many AI applications depend on **access to current operational context**. The challenge is not that organizations lack data. In many cases, they already have it. The challenge is that the data is spread across multiple systems, each with its own copy, latency profile, and synchronization process.

"AI is forcing organizations to confront a broader question: How many copies of the same data are actually necessary? The answer increasingly appears to be fewer than most architectures maintain today."

As a result, AI is forcing organizations to confront a broader question: How many copies of the same data are actually necessary? The answer increasingly appears to be fewer than most architectures maintain today. As expectations around freshness rise, reducing unnecessary data movement becomes just as important as

What's next

The database industry spent decades solving storage. Databases became more reliable, storage became cheaper, and infrastructure became dramatically easier to operate. The next challenge is not where data lives, but how easily it can be shared across systems without introducing additional pipelines, copies, and synchronization overhead. Increasingly, the goal is not simply moving data faster. It is reducing unnecessary movement altogether.

Postgres has a habit of outlasting predictions about its replacement. For years, members of the community have joked that every year is “the year of Postgres.” The joke works because it keeps turning out to be true.

Three decades after its creation, **Postgres continues to adapt** to new workloads, new architectural patterns, and new ways of building applications.

That longevity is not an accident. Enterprises continue to rely on Postgres because it provides a stable and trusted foundation for operational data. While that foundation is unlikely to change, the scope of what organizations expect Postgres to do will continue to expand.

As new workloads continue to emerge, much of the innovation will come through extensions that expand Postgres’s capabilities without sacrificing the stability that made it successful. In that sense, the future of Postgres may not be about reinventing the database itself, but continuously expanding what can be built on top of it.

TNS



Craig Kerstiens leads product and engineering efforts for the Snowflake Postgres team, helping drive Snowflake’s Postgres compatibility and the broader database ecosystem. He is responsible for enabling enterprises to adopt Postgres at scale and integrating open source database innovation into...

[Read more from Craig Kerstiens →](#)

TRENDING STORIES

1. The database storage problem is solved. Here's what comes next.
2. When your data model is the bottleneck: lessons from Medium's feature store
3. Why MotherDuck refuses to fork DuckDB
4. Why the need for humans won't disappear in the age of autonomous databases
5. For years, Apache Cassandra handed this work to your team — 6.0 takes it back

INSIGHTS FROM OUR SPONSOR



Snowflake enables every organization to mobilize their data with Snowflake's Data Cloud. Customers use the Data Cloud to unite siloed data, discover and securely share data, power data applications, and execute diverse AI/ML and analytic workloads across multiple clouds and geographies.

[Learn More →](#)

TNS DAILY NEWSLETTER

most recent TNS articles in your inbox each day.

svo1975pvn1982@gmail.com

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter

THE NEW STACK

[About / Contact](#)
[Sponsors](#)
[Advertise With Us](#)
[Contributions](#)



Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

[Frontend Developer Roadmap](#)
[Backend Developer Roadmap](#)
[Devops Roadmap](#)

FOLLOW TNS



© The New Stack 2026

[Disclosures](#)
[Terms of Use](#)
[Advertising Terms & Conditions](#)
[Privacy Policy](#)
[Cookie Policy](#)