


ИИ / МОДЕЛИ ИИ / СТРАТЕГИЯ ИИ / БЕЗОПАСНОСТЬ

Правительство США только что сообщило OpenAI, кому разрешено использовать следующую модель GPT 5.6

Спустя несколько недель после введения ограничений Anthropic в отношении Claude Белый дом приказал OpenAI ограничить доступ к GPT-5.6 для проверенных партнеров, сославшись на соображения национальной безопасности.

26 июня 2026 года, 14:26 [Адриан Бриджуотер](#)



Ангга Дви Баю Сапутра

Выпуск ведущих базовых моделей искусственного интеллекта достиг переломного момента. Всего через две недели после того, как **Anthropic получила распоряжение** от правительства США приостановить доступ к Fable 5 и Mythos 5, Белый дом на этой неделе обязал OpenAI ввести ограничения на выпуск будущей модели GPT 5.6 из соображений кибербезопасности.

Согласно новой директиве, GPT 5.6 от OpenAI может быть предоставлен только ограниченному числу партнеров, одобренных правительством, из-за

В статье, опубликованной в среду на *The Information*, упоминается сессия вопросов и ответов с участием сотрудников, организованная генеральным директором OpenAI **Сэмом Альтманом**. В **посте на X**, основанном на этой статье, объясняется, что федеральное правительство будет «выдавать разрешения на доступ клиентам по очереди в течение этого периода предварительного тестирования» GPT 5.6.

Заккрытие двери подливает масла в огонь

Коллин Хог-Спирс, старший директор и выдающийся эксперт в области технологий в **Black Duck**, рассказывает *The New Stack*, что закрытие доступа к API не снижает возможности, а смещает спрос в те регионы, где правительство США не имеет юрисдикции.

разработчики, поэтому контроль доступа к закрытой модели подталкивает серьезные команды к поиску альтернативы, которую они могли бы использовать самостоятельно, — говорит Хог-Спирс. — Каждое одобрение со стороны клиента делает эту альтернативу более привлекательной, а не наоборот, так что цель обеспечения безопасности и реакция рынка теперь противоречат друг другу».

“...the security goal and the market reaction now run against each other.”

Hogue-Spears advises that the absence of federal regulation is the “core problem for builders,” who need a fixed rule to plan against rather than discretionary action applied on a model-by-model basis. At the same time, he says teams are navigating a patchwork of state AI laws with conflicting definitions and deadlines, which, in his view, is a “result of the lack of leadership” from Washington.

Teams are navigating a patchwork of state AI laws

While it's true that no formalized federal regulatory framework exists to govern and oversee the ongoing release of new AI models, President **Donald Trump signed an executive order** on June 2 asking AI companies to submit new model releases 30 days before release.

TRENDING STORIES

1. **OpenClaw and Hermes agree on what an agent is. They disagree on what controls it.**
2. **MCP gets its missing enterprise authorization layer**
3. **Developers are now validating code they didn't write — and may not understand**
4. **Can DNS become the basis for AI agent identity?**
5. **AI agents need infrastructure: Why Europe's regional cloud strategy matters**

capabilities make the nation stronger.

“We’ve made clear to the U.S. government that this is not our preferred long-term model, and will work with them and others in industry to achieve a more sustainable approach for future releases.” —Sam Altman.

“But [they] also introduce new national security considerations that require coordinated action across executive departments and agencies (agencies) and components. As these capabilities evolve, my Administration will continue to work closely with industry to ensure that the best and most secure technology is deployed rapidly to confront any and all threats to our country,” details the executive order signed by Trump.

In a snippet from the OpenAI CEO’s memo seen by *The Information*, Altman is reported to have been less than happy: “We’ve made clear to the U.S. government that this is not our preferred long-term model, and will work with them and others in industry to achieve a more sustainable approach for future releases.”

What software developers think of government directives

The software engineering community has, unsurprisingly, been unimpressed. CIO at Interconnected Capital and ex-GitHub and Obama White House/Commerce Dept executive **Kevin S. Xu** **posted on X** saying, “The government will be approving access to GPT-5.6 customer by customer? [That makes China’s model pre-release registration scheme look almost laissez-faire.”

Echoing that sentiment, [Charming-Author4877](#) writes on Reddit, “By the time they release GPT 5.6 we’ll hopefully have the next GLM, Qwen, Deepseek or Kimi that beats it. [The] US is within months of losing the lead in AI.”

Redditor [Punter1965](#) agrees with the general sentiments expressed thus far. He has suggested that all this will do is drive companies to choose open-source or Chinese models rather than deal with these headaches.

“Basically, Trump’s spat with Anthropic and its spillover will only diminish America’s lead,” they noted. “Unfortunately, by the time this gets worked out (either by

With developers joking that next, we'll see the "government is approving GPT-5.6 usage prompt-by-prompt", the mood is not upbeat.

"My concern is the long-term impact: if access to cutting-edge AI increasingly depends on where you live or which organization you work for, we risk moving away from the open innovation model." – Sylwia Laskowska

The Chrome-style release honeymoon period is over

Gdańsk, Poland-based senior software engineer **Sylwia Laskowska** tells *The New Stack* that until now, new AI models have been released much like new versions of Chrome or React, i.e developers could start experimenting with them almost immediately.

"If access to cutting-edge AI increasingly depends on where you live or which organization you work for, we risk moving away from the open innovation model."

"If frontier AI models are now becoming strategic assets, similar to advanced semiconductors, that's a fundamental shift," Laskowska says. "For the average developer, this may not make much difference at first. We'll probably just receive more stable, thoroughly tested releases. My concern is the long-term impact: if access to cutting-edge AI increasingly depends on where you live or which organization you work for, we risk moving away from the open innovation model."

CISO at Sumo Logic, **Jeremy Powell**, tells *The New Stack* that the real story here is about capability, rather than politics. He thinks that, from a security perspective, when a lab and the government both decide a model is too powerful to release widely, that tells us the offensive bar just moved.

doesn't undo the capability."

"For developers, that ability to connect up different issues between components to achieve an exploit would mean that they have to update those components more rapidly," Powell says. "Restricting access to frontier AI models slows the spread — it doesn't undo the capability. At the same time, that improvement in performance will make a difference to developers around how they build software and create new agents that can handle more complex tasks."

TNS



Adrian Bridgwater is a technology journalist with three decades of press experience. He has an extensive background in communications, starting in print media, newspapers and also television. Primarily working as an analysis writer dedicated to a software application development 'beat',...

[Read more from Adrian Bridgwater →](#)

TNS owner Insight Partners is an investor in: OpenAI, Anthropic.

SHARE THIS STORY



TRENDING STORIES

1. OpenClaw and Hermes agree on what an agent is. They disagree on what controls it.
2. MCP gets its missing enterprise authorization layer
3. Developers are now validating code they didn't write — and may not understand
4. Can DNS become the basis for AI agent identity?
5. AI agents need infrastructure: Why Europe's regional cloud strategy matters



How to Eliminate Training-Serving Skew With a Unified Real-Time Streaming ML Pipeline (2026 Guide)

23 June 2026

Real-Time Hyper-Personalization in 2026: Architecture Guide

23 June 2026

Meeting Data (and Analytics) Engineers Where They Are: Introducing the dbt Adapter for Confluent Cloud

18 June 2026



NUMA Explained: Why Memory Distance Slows Your VMs

23 June 2026

Docker-in-Docker Security Risks and How to Escape Them

17 June 2026

Kubernetes User Namespaces Don't Fix the Shared Kernel

4 June 2026



Multicloud HIPAA compliance for healthcare: Dynatrace now supports AWS, Azure, and GCP

26 June 2026

Beyond LLM-as-a-judge: Establishing LLM evaluations as a foundation for trustworthy agentic AI systems

June 2026

TNS DAILY NEWSLETTER

Receive a free roundup of the
most recent TNS articles in
your inbox each day.

svo1975pvn1982@gmail.com

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking

ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions

 roadmap.sh

Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

FOLLOW TNS



© The New Stack 2026

Disclosures
Terms of Use
Advertising Terms & Conditions
Privacy Policy
Cookie Policy

